



Lecture 1: Introduction

Xin Chen

Assistant Professor

University of New Mexico

What will we learn?

- The basic algorithms in **traditional learning** and their implementations in Python.
- Basic concepts in **supervised learning**, **unsupervised learning** and **reinforcement learning**.
- Applications of traditional learning algorithms.
- The ML libraries [scikit-learn](#) and [PyTorch](#).
- Building and training **artificial neural networks** in PyTorch.
- Applications of artificial neural networks.

How to learn this course?

- **Taking every class.** Since we have a comprehensive course content, each lecture is important.
- **Taking notes in every class.** The slides will not include everything, please take notes of the content that is not contained in the slides and textbook.
- **Reading the textbook.** The slides are not a textbook, you should never expect to understand the course content by only reading the slides.
- **Reading online materials.** Some online references will be provided in the slides, however, you are encouraged to proactively find online materials using the skills and knowledge learned in this course.
- **Reviewing after every class.** Please read the corresponding chapter(s) and the lecture slides after every class.
- **Arranging regular team discussion.** Please discuss with your teammates regularly.

Somethings you need to know

- During a class, **never use laptop(s), cell phone(s) or other things** that distract you from the class.
- This is not a basic programming course, **we do not help with debugging**.
- Please let me know in advance if you will be late (after 10 min) to a class. Please do not be late to a class frequently.
- If you meet any difficulty in course study, please reach out to the instructor or TA **at the first convenient time**.
- **AI tools**, including any LLMs, are **forbidden** in this course.

What is Machine Learning (ML)?

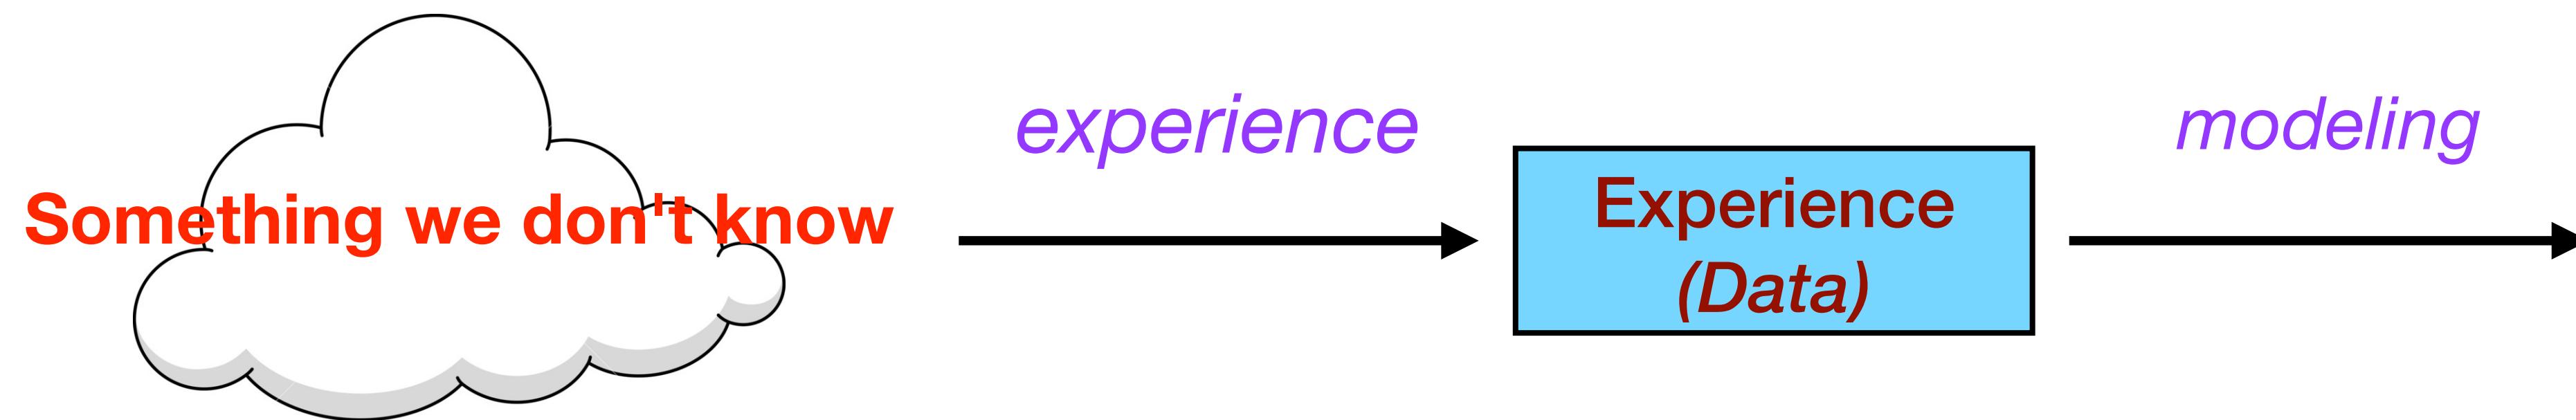
"Learning is any process by which a system improves performance from experience."

- Herbert Simon

Machine learning is the study of algorithms that

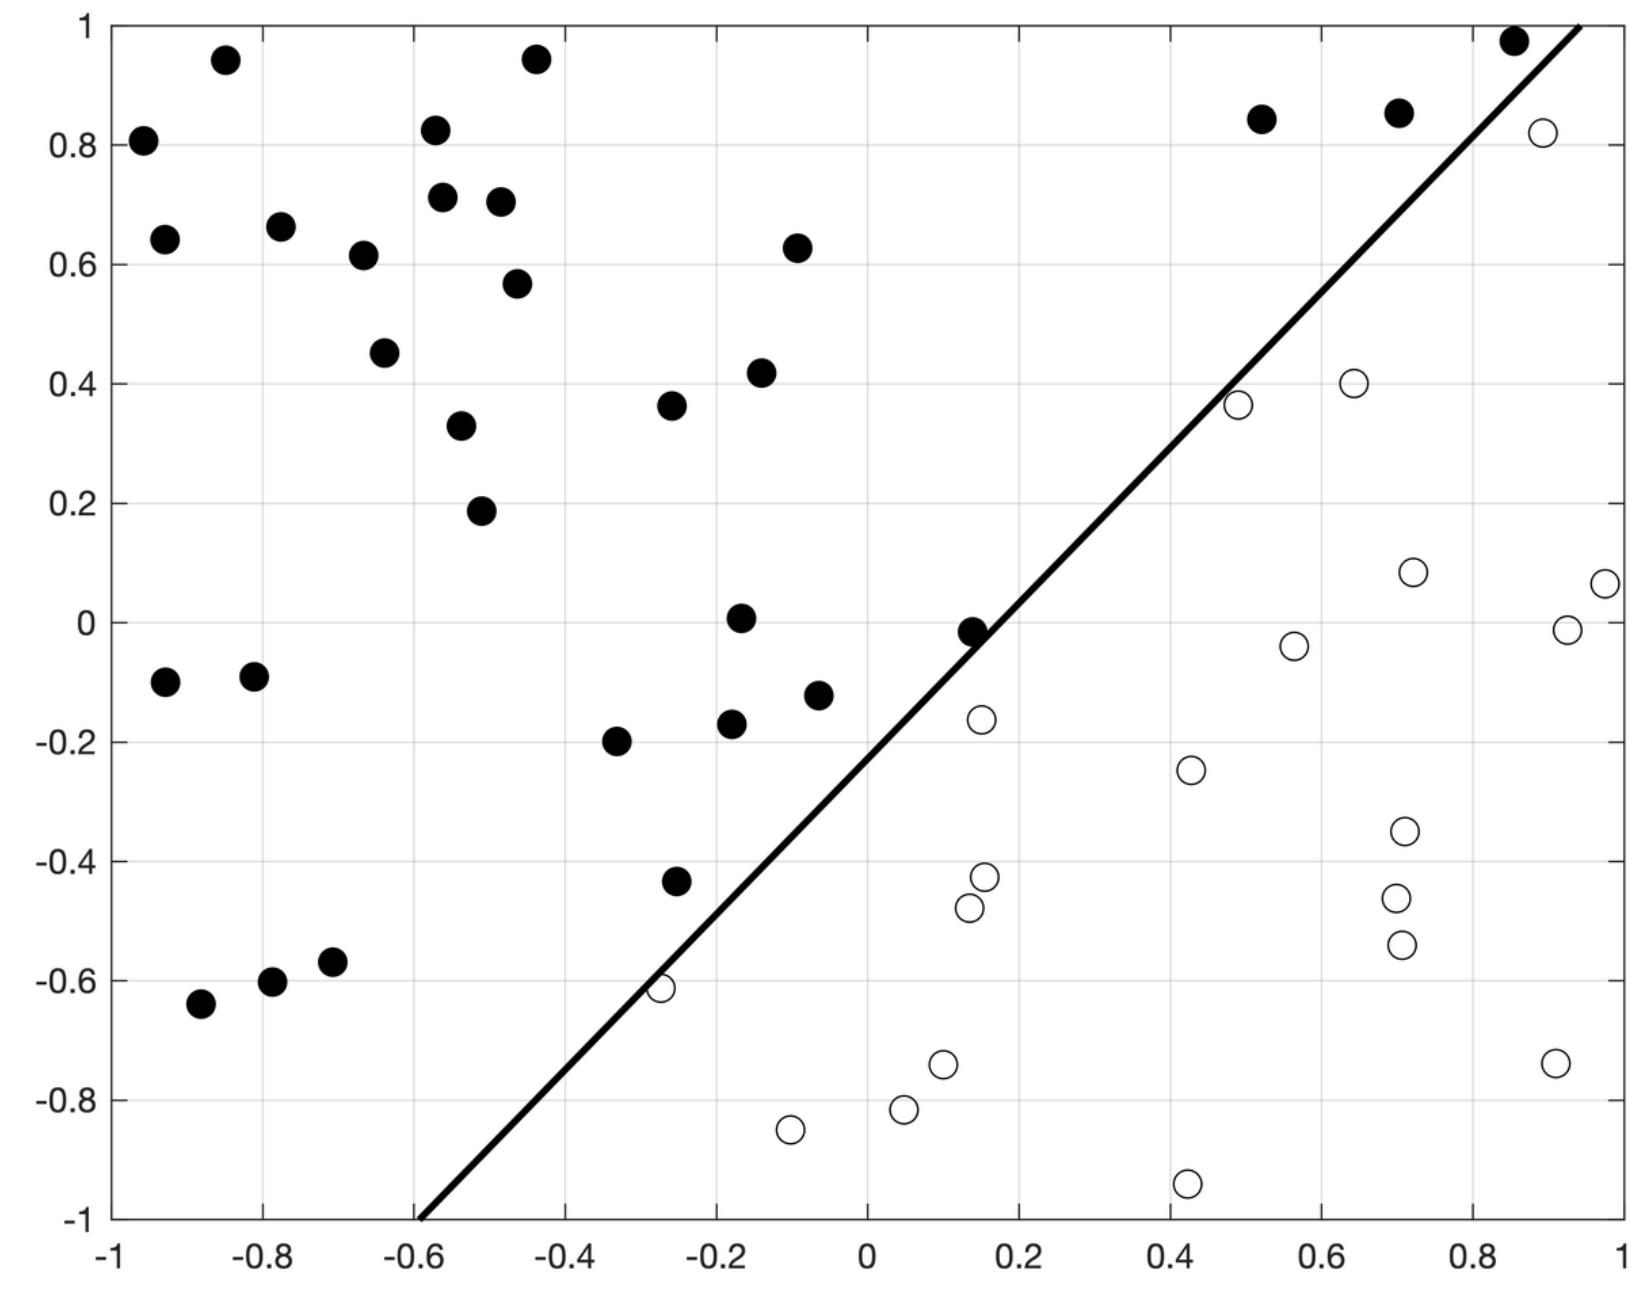
- improve the performance P ,
- at some task T ,
- and with experience E .

- Tom Mitchell

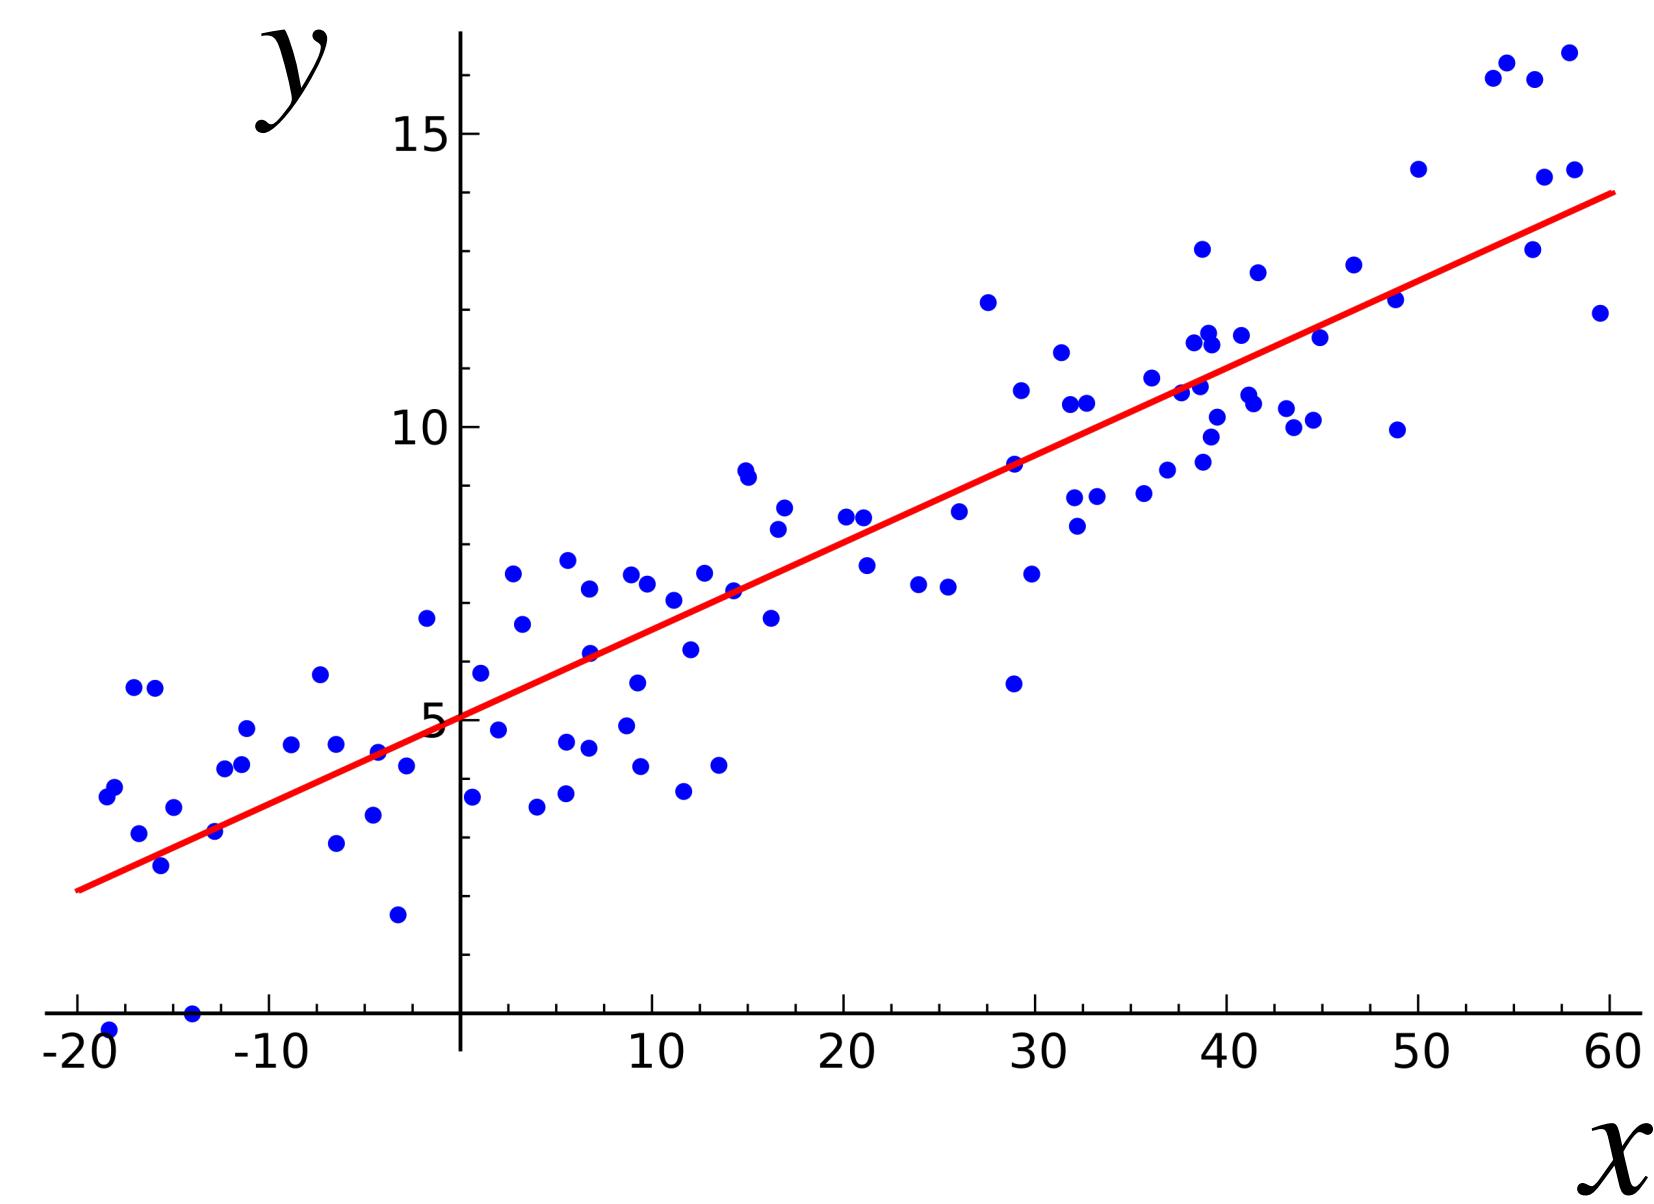


$$\begin{aligned}x_{k+1} &= g(x_k) \\dx/dt &= f(x, t) \\\bar{x} &= h(\bar{y})\end{aligned}$$

Example - Tasks of Supervised Learning



Classification: Finding a *hyperplane* that distinguishes the black and white nodes.



Regression: Finding a *function* that describes the relation between the variables x and y .

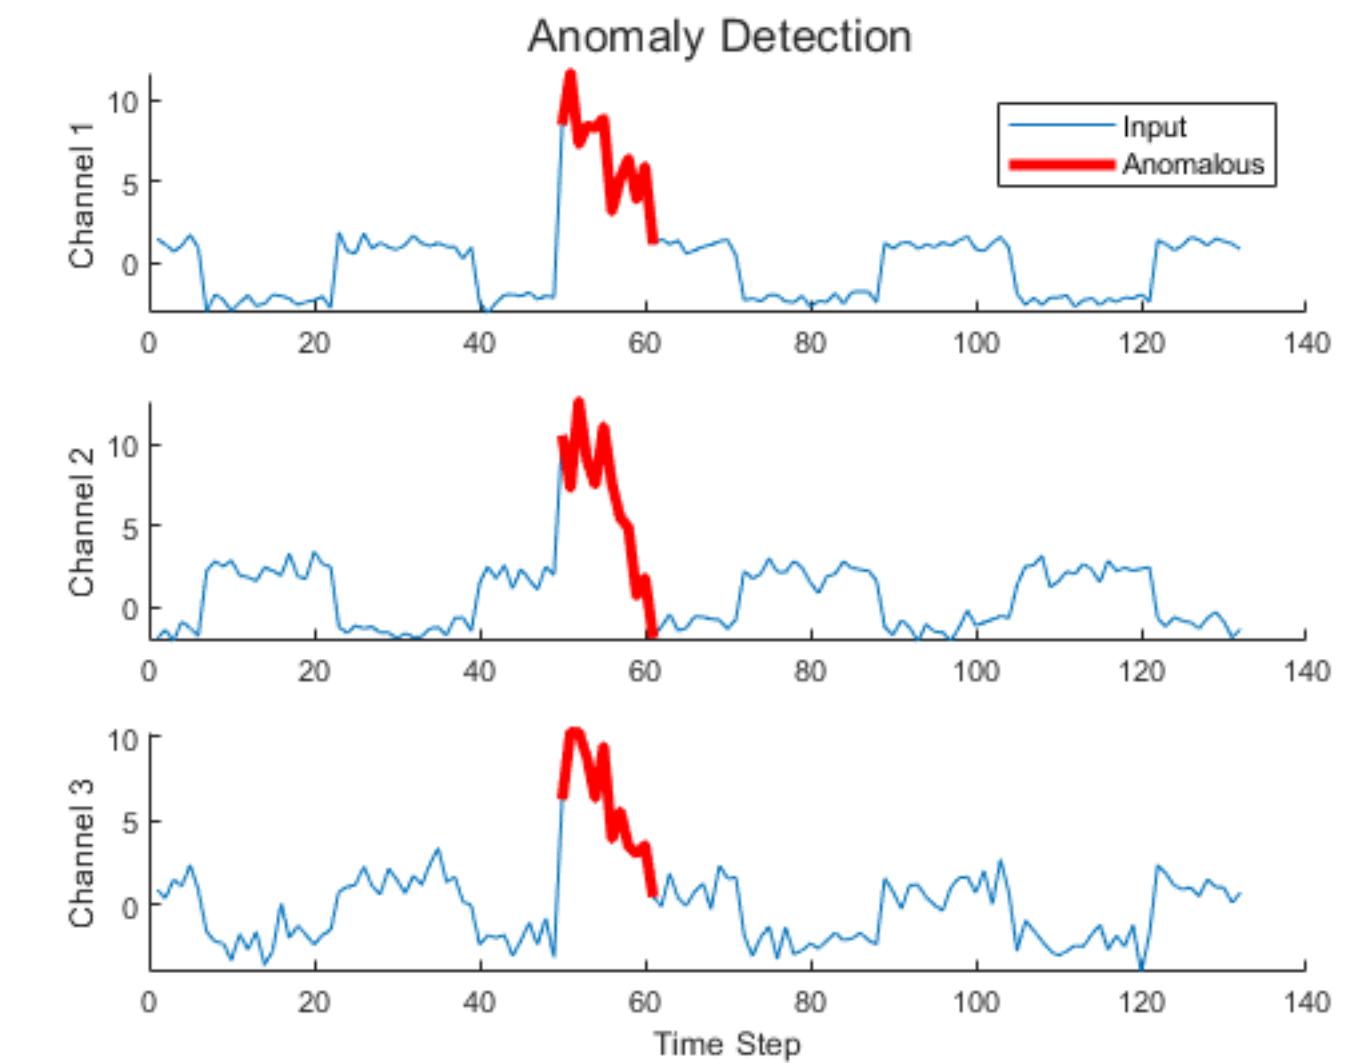
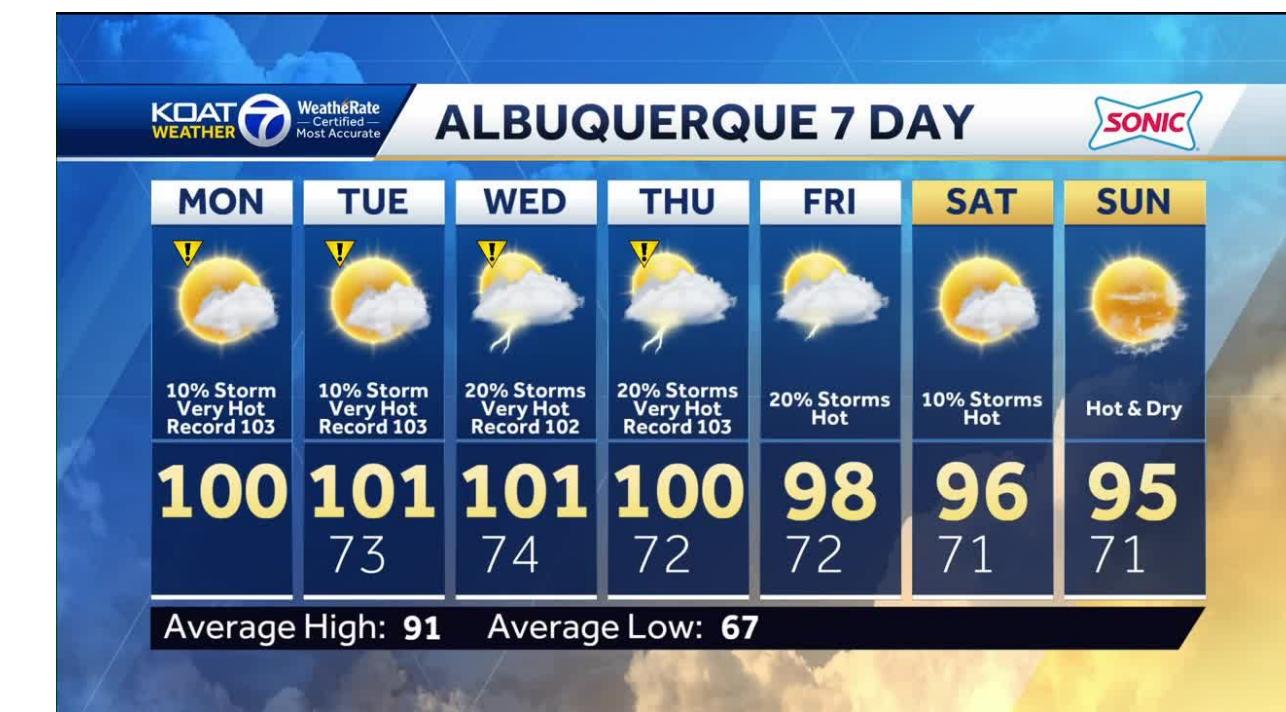
Advanced ML Tasks

Anomaly Detection

Pattern Recognition



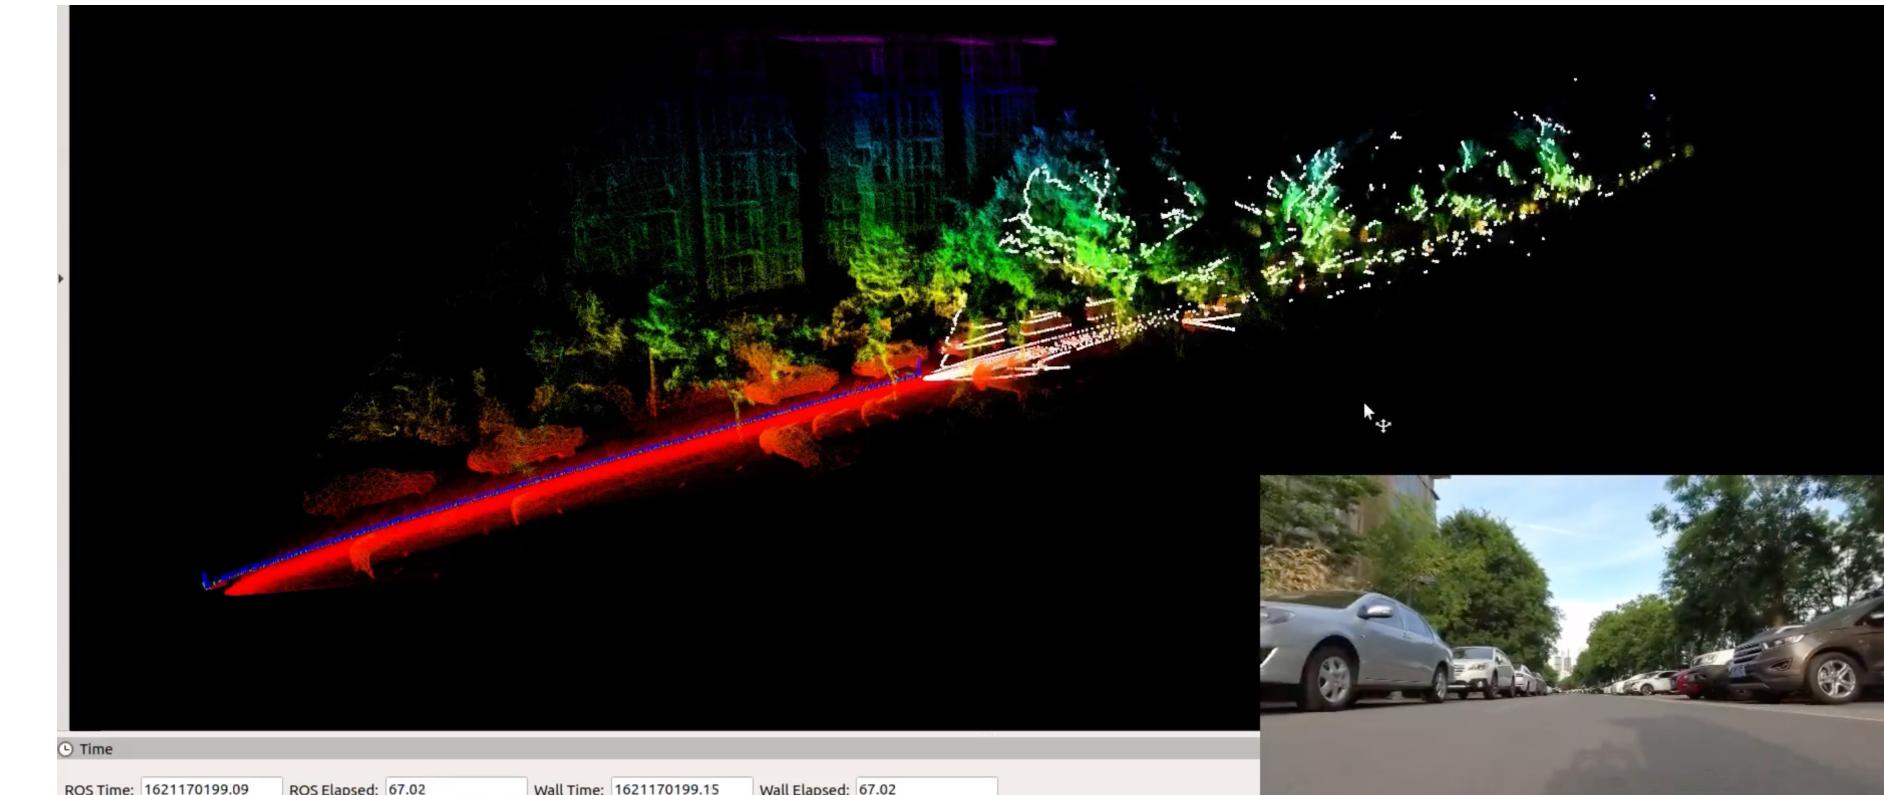
Weather Forecast



Voice Translation

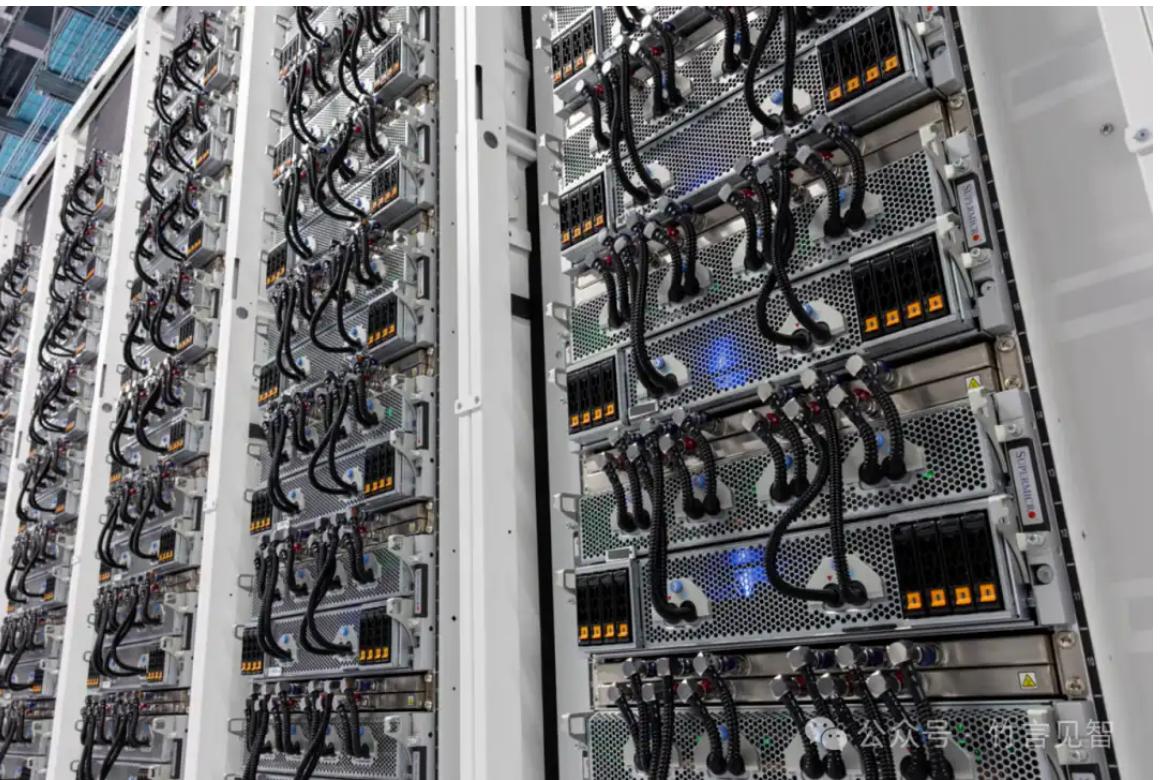


Motion Planning



Hardware for ML

Training ML Models



Using ML Models



AI/ML in Embedded Systems

Hardware

Raspberry Pi



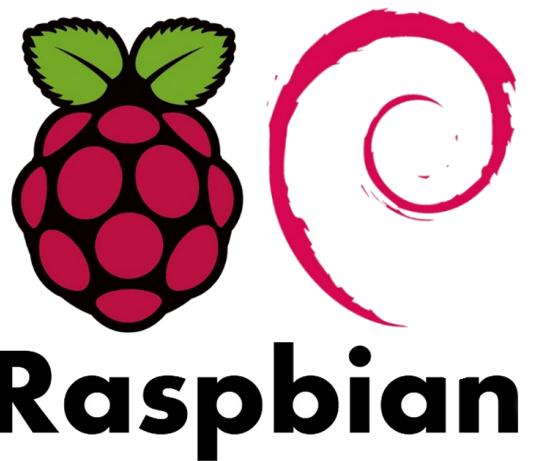
Nvidia Jetson



D-Robotics RDK



Software

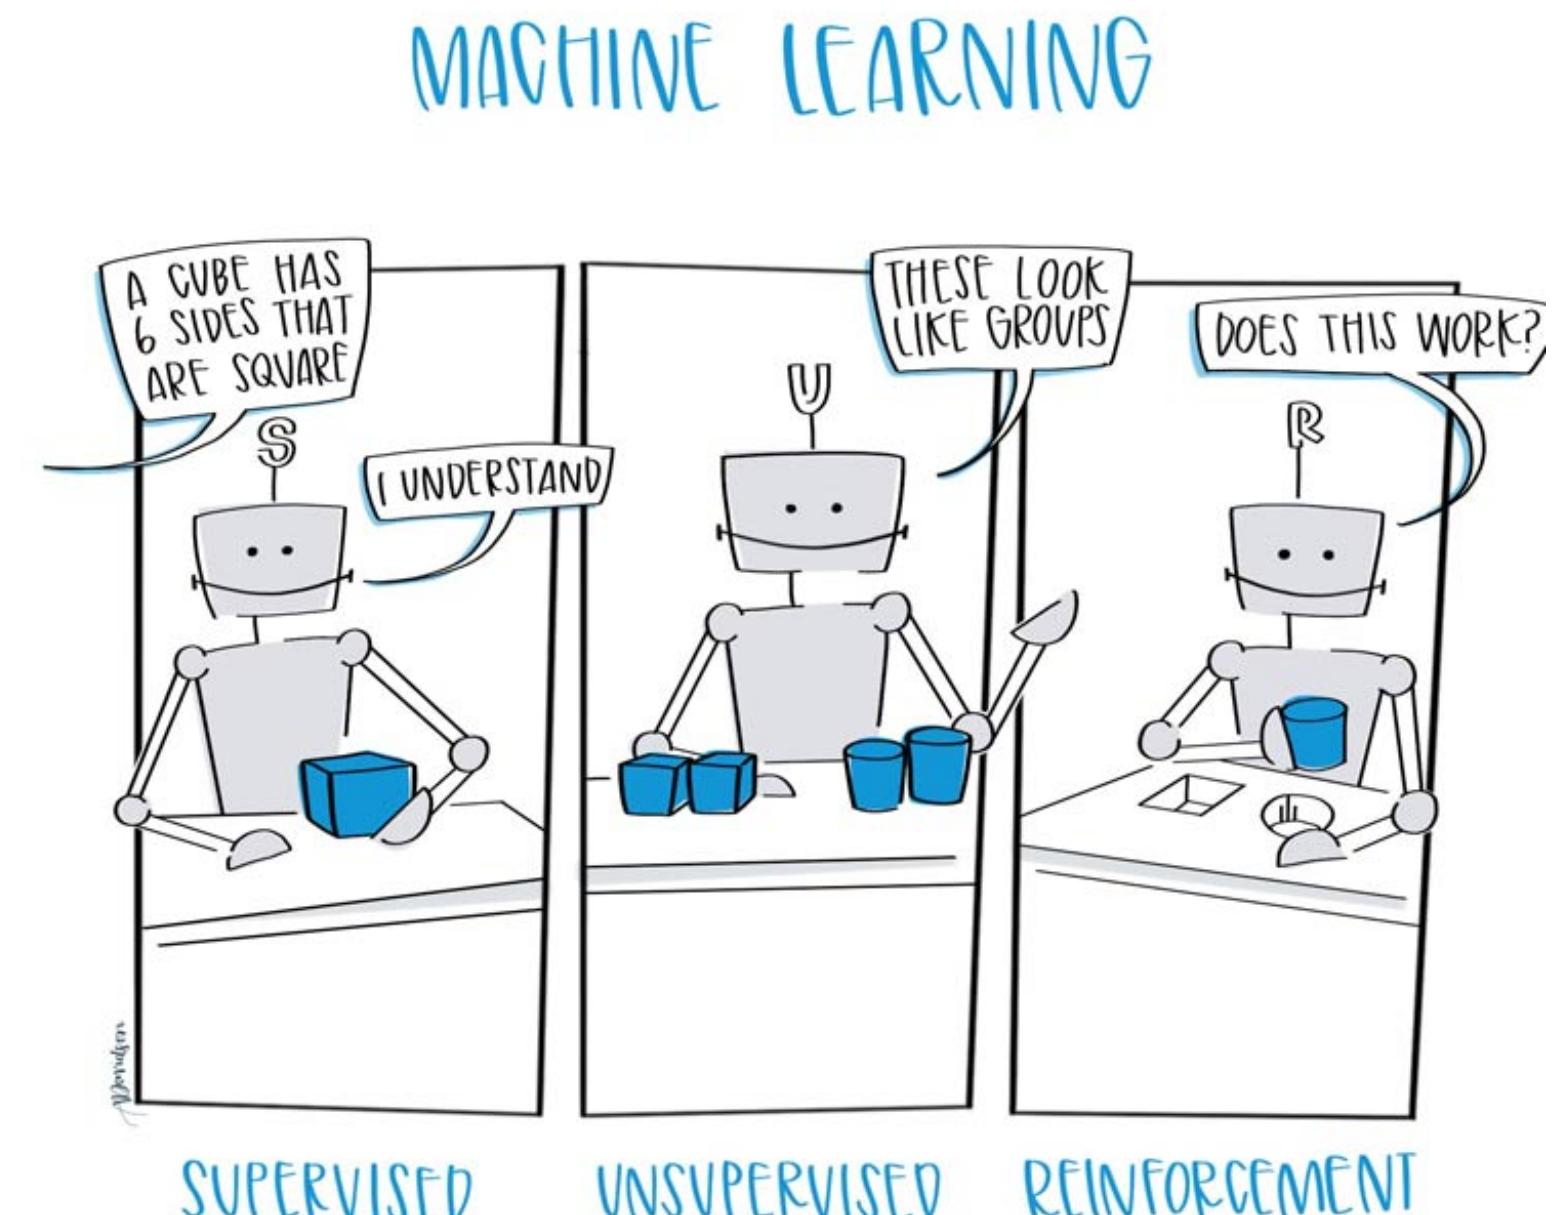


Application



Types of Machine Learning

- **Supervised Learning:**
 - ▶ *Classification* and *regression*.
- **Unsupervised Learning:**
 - ▶ *Clustering*, *dimensionality reduction*, *autoencoders*.
- **Reinforcement Learning:**
 - ▶ *Finding a policy based on finite interactions with the environment.*



 ceralytics

Supervised Learning

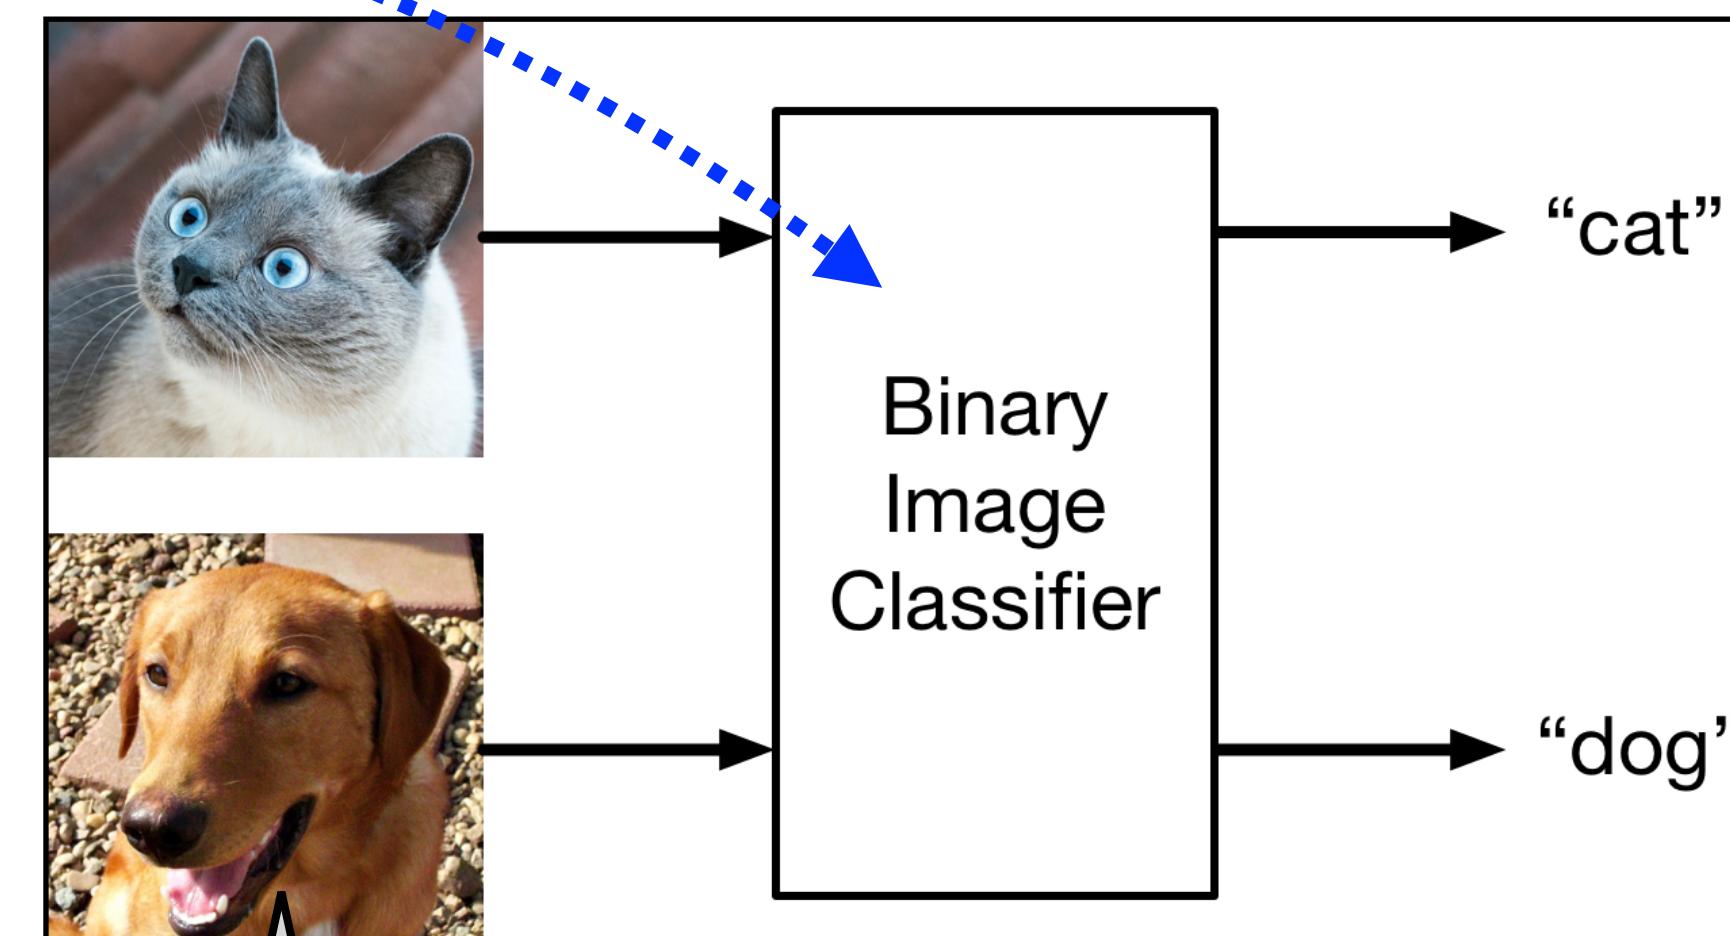
Training Data/Samples

Each sample is correctly labeled. A training data set might not need to be large but is expected to be comprehensive.



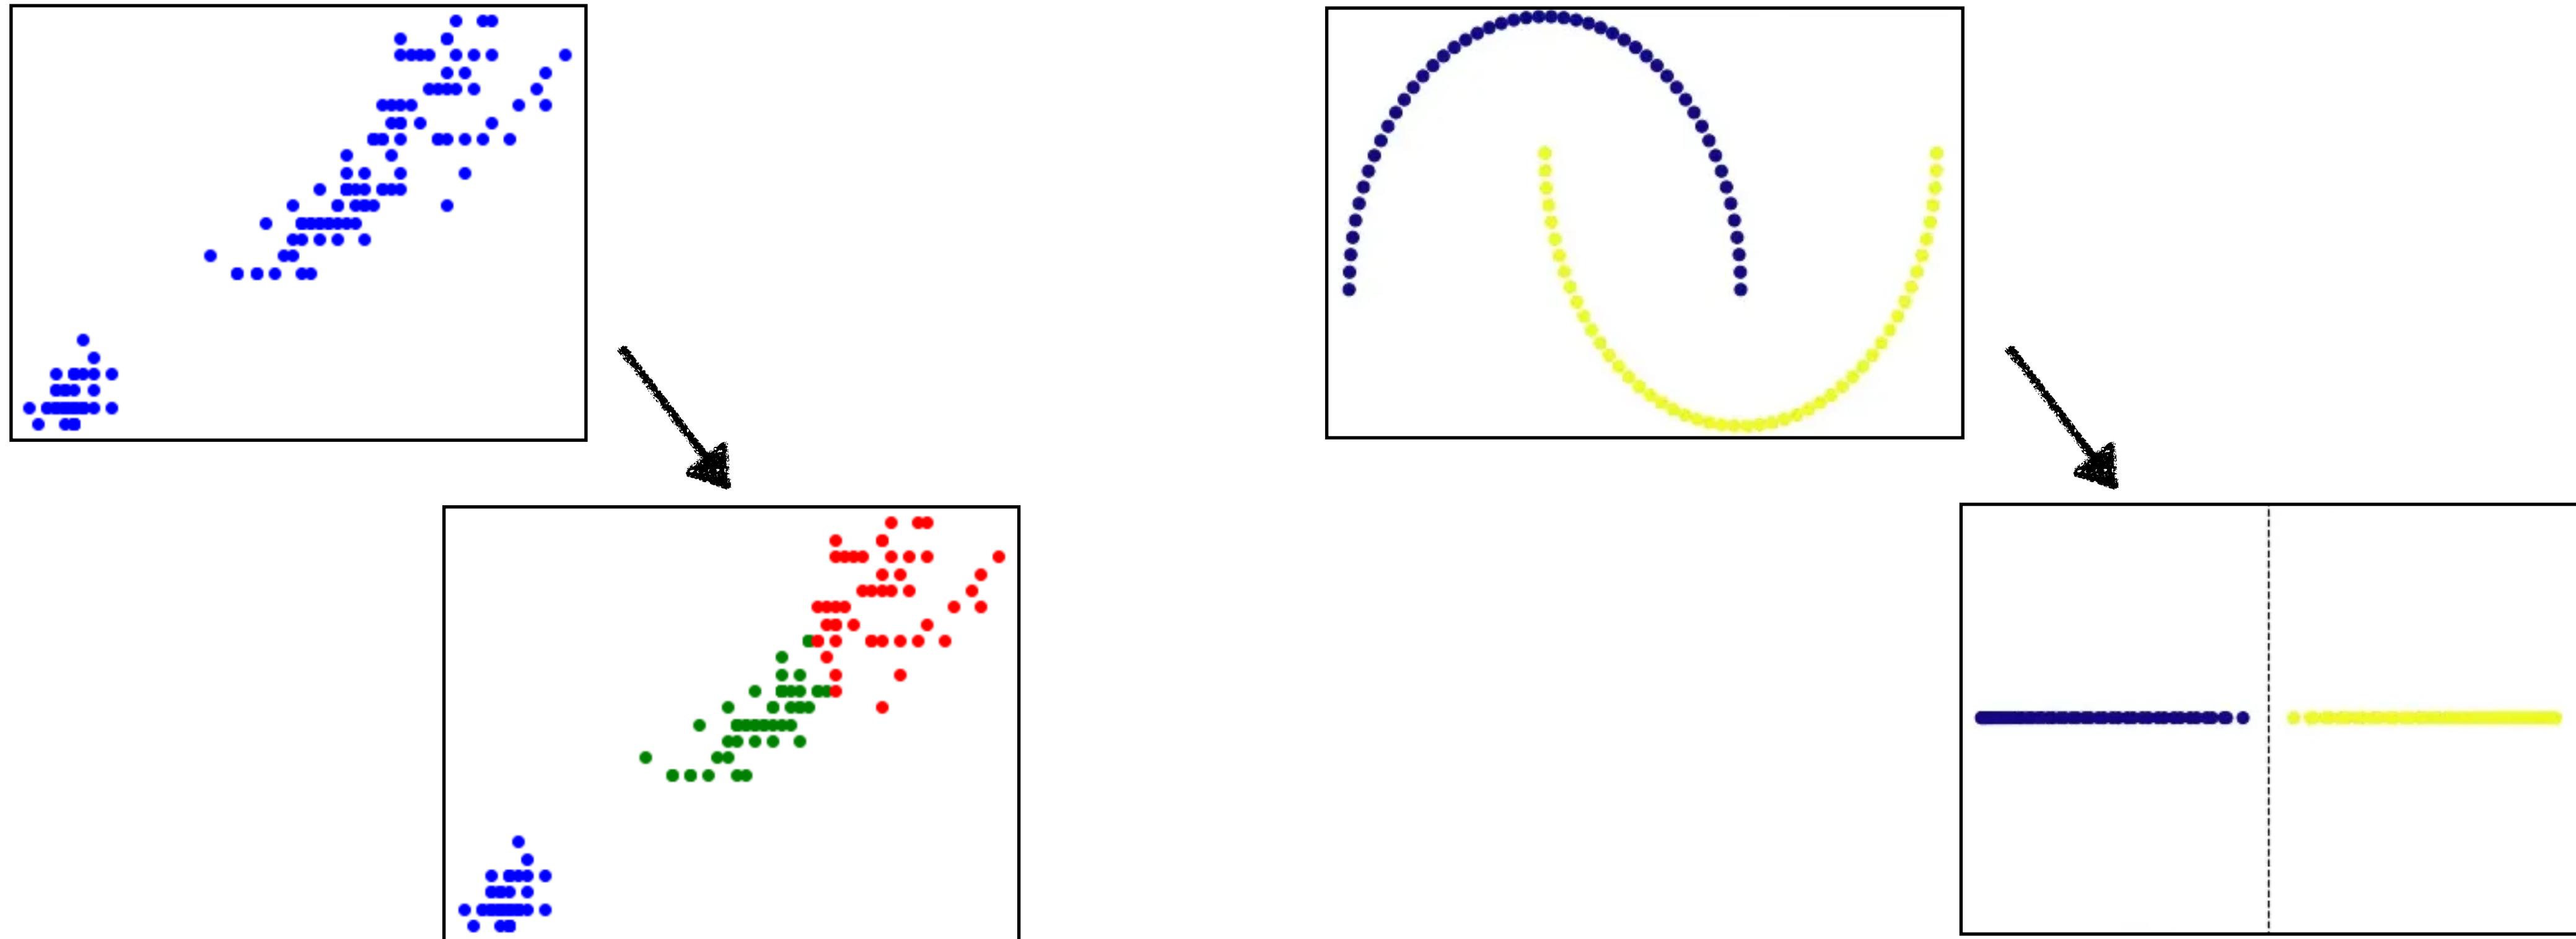
Training an ML Model

Finding the unknown parameters in a given model to minimize its prediction error on the training samples.



Samples used to evaluate the performance of the ML model. These samples should not be in the training data set.

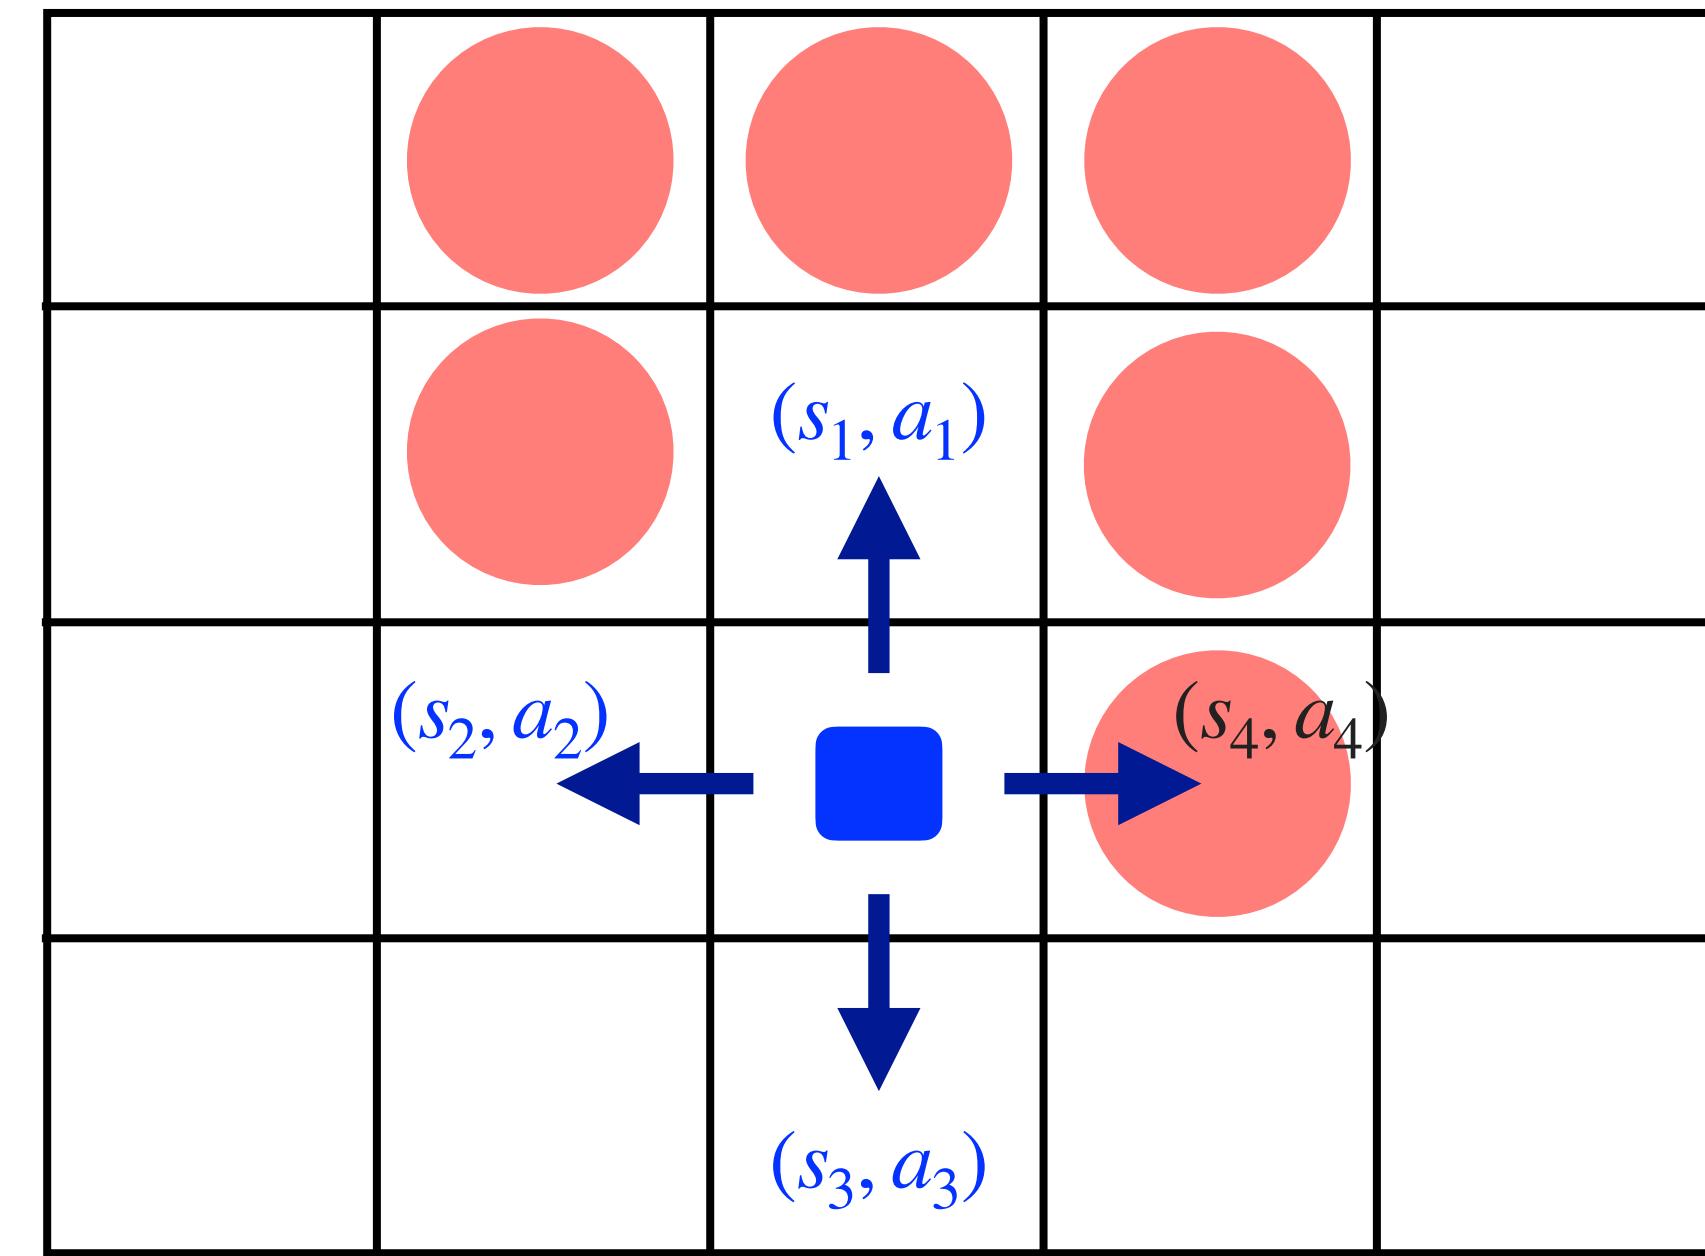
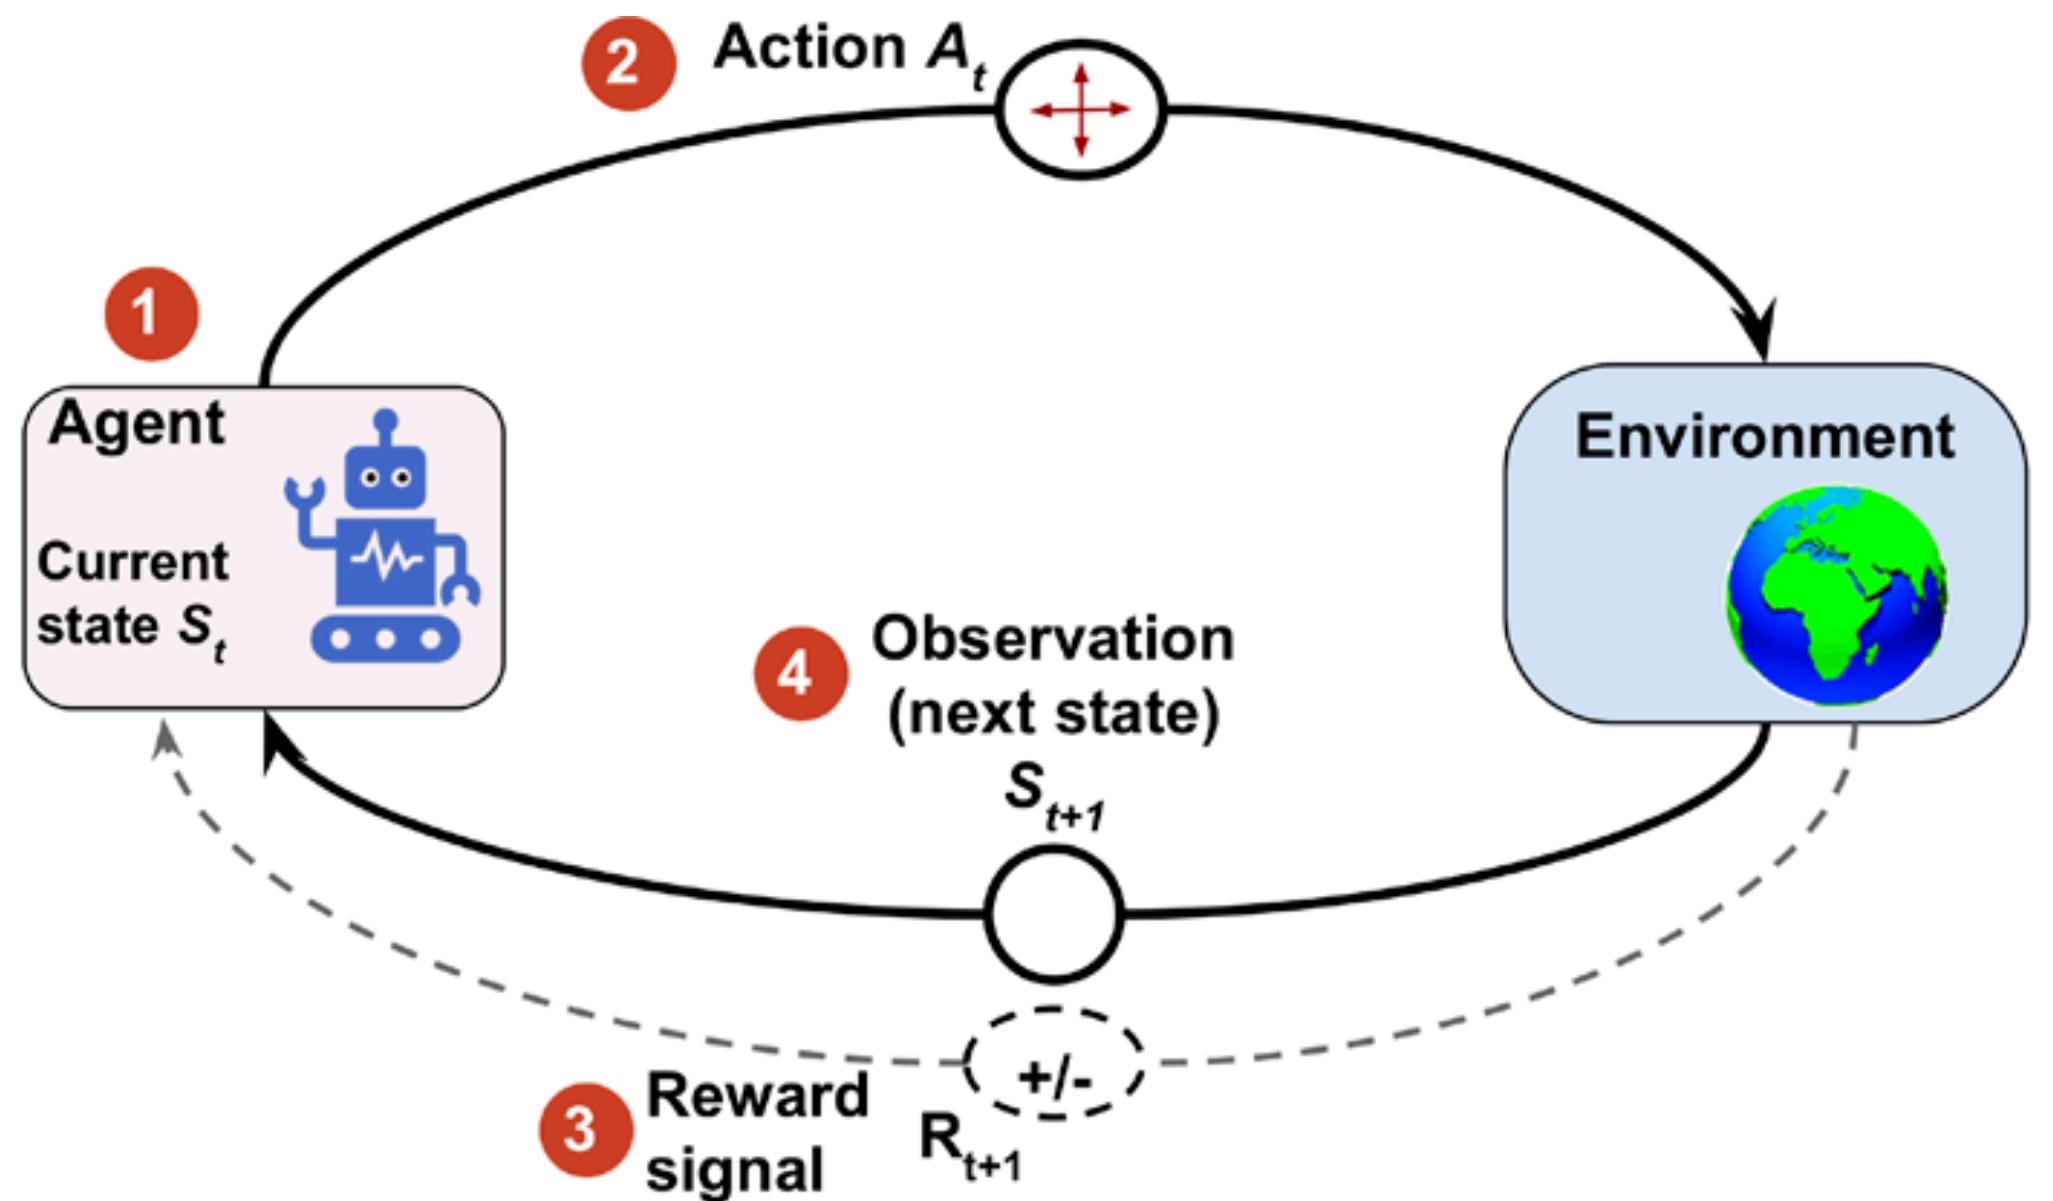
Unsupervised Learning



Clustering: Partitioning a set of (unlabeled) samples into homogeneous subsets.

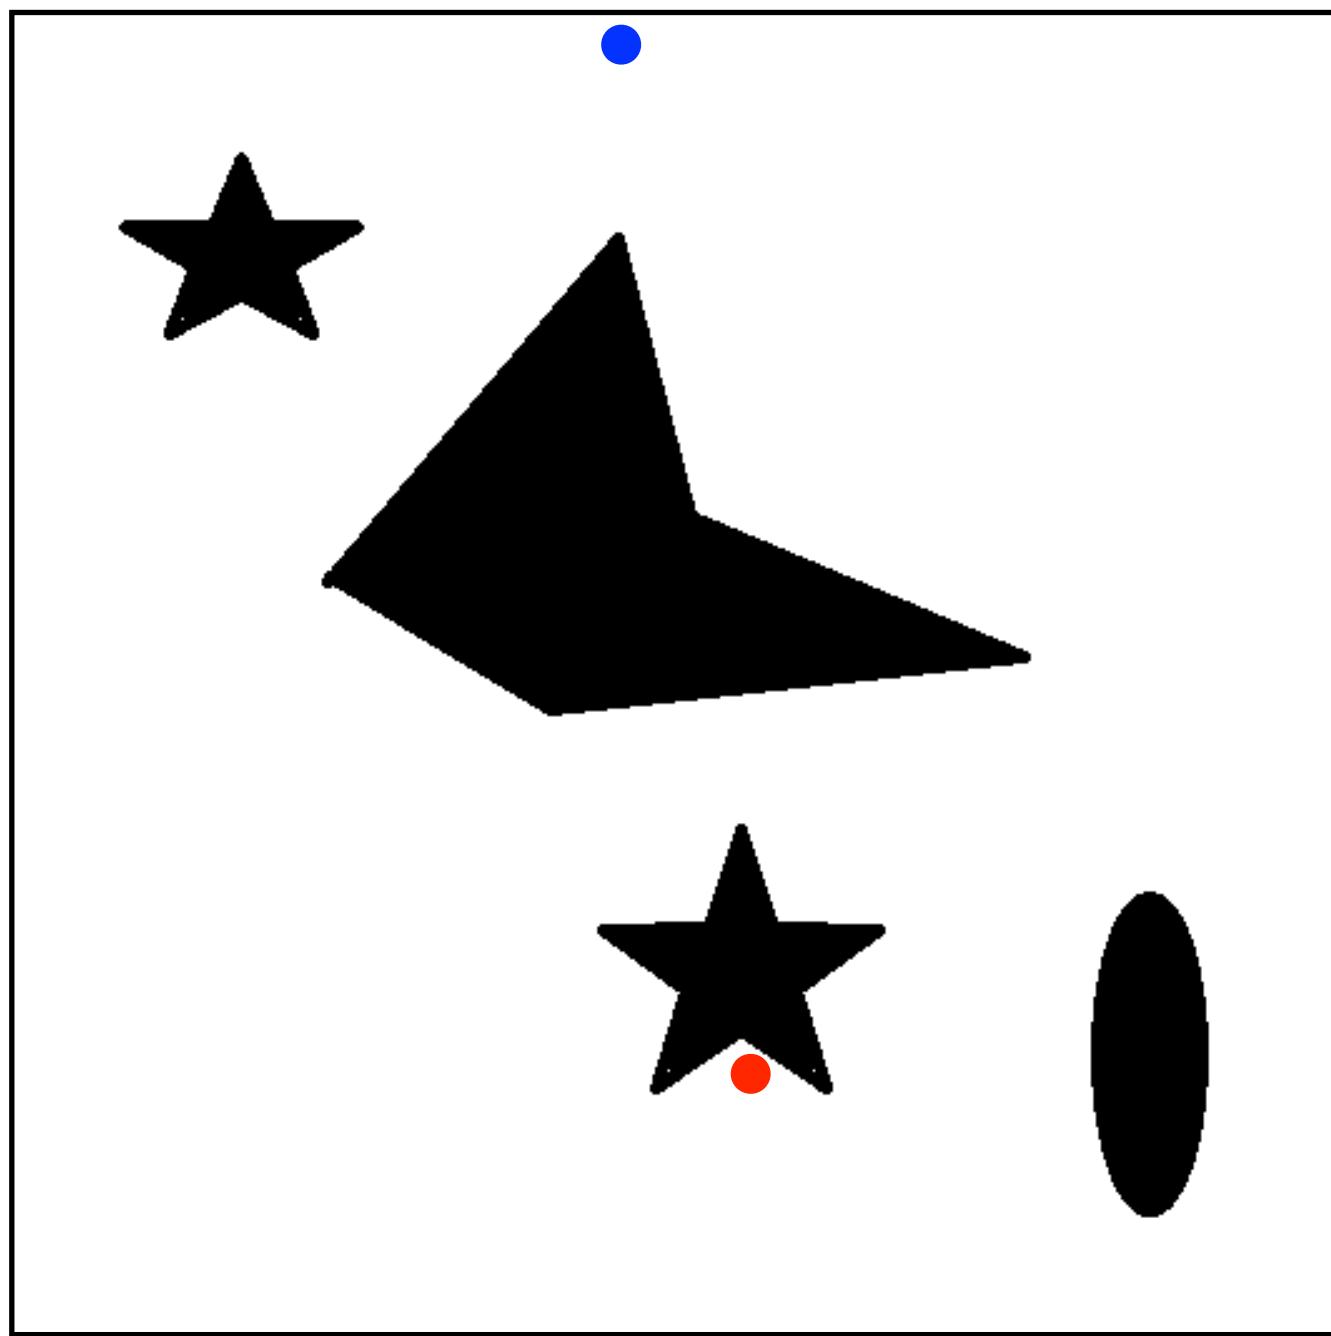
Dimensionality Reduction: Transforming an initial representation of samples into a lower-dimensional representation while preserving some properties of the initial representation.

Reinforcement Learning

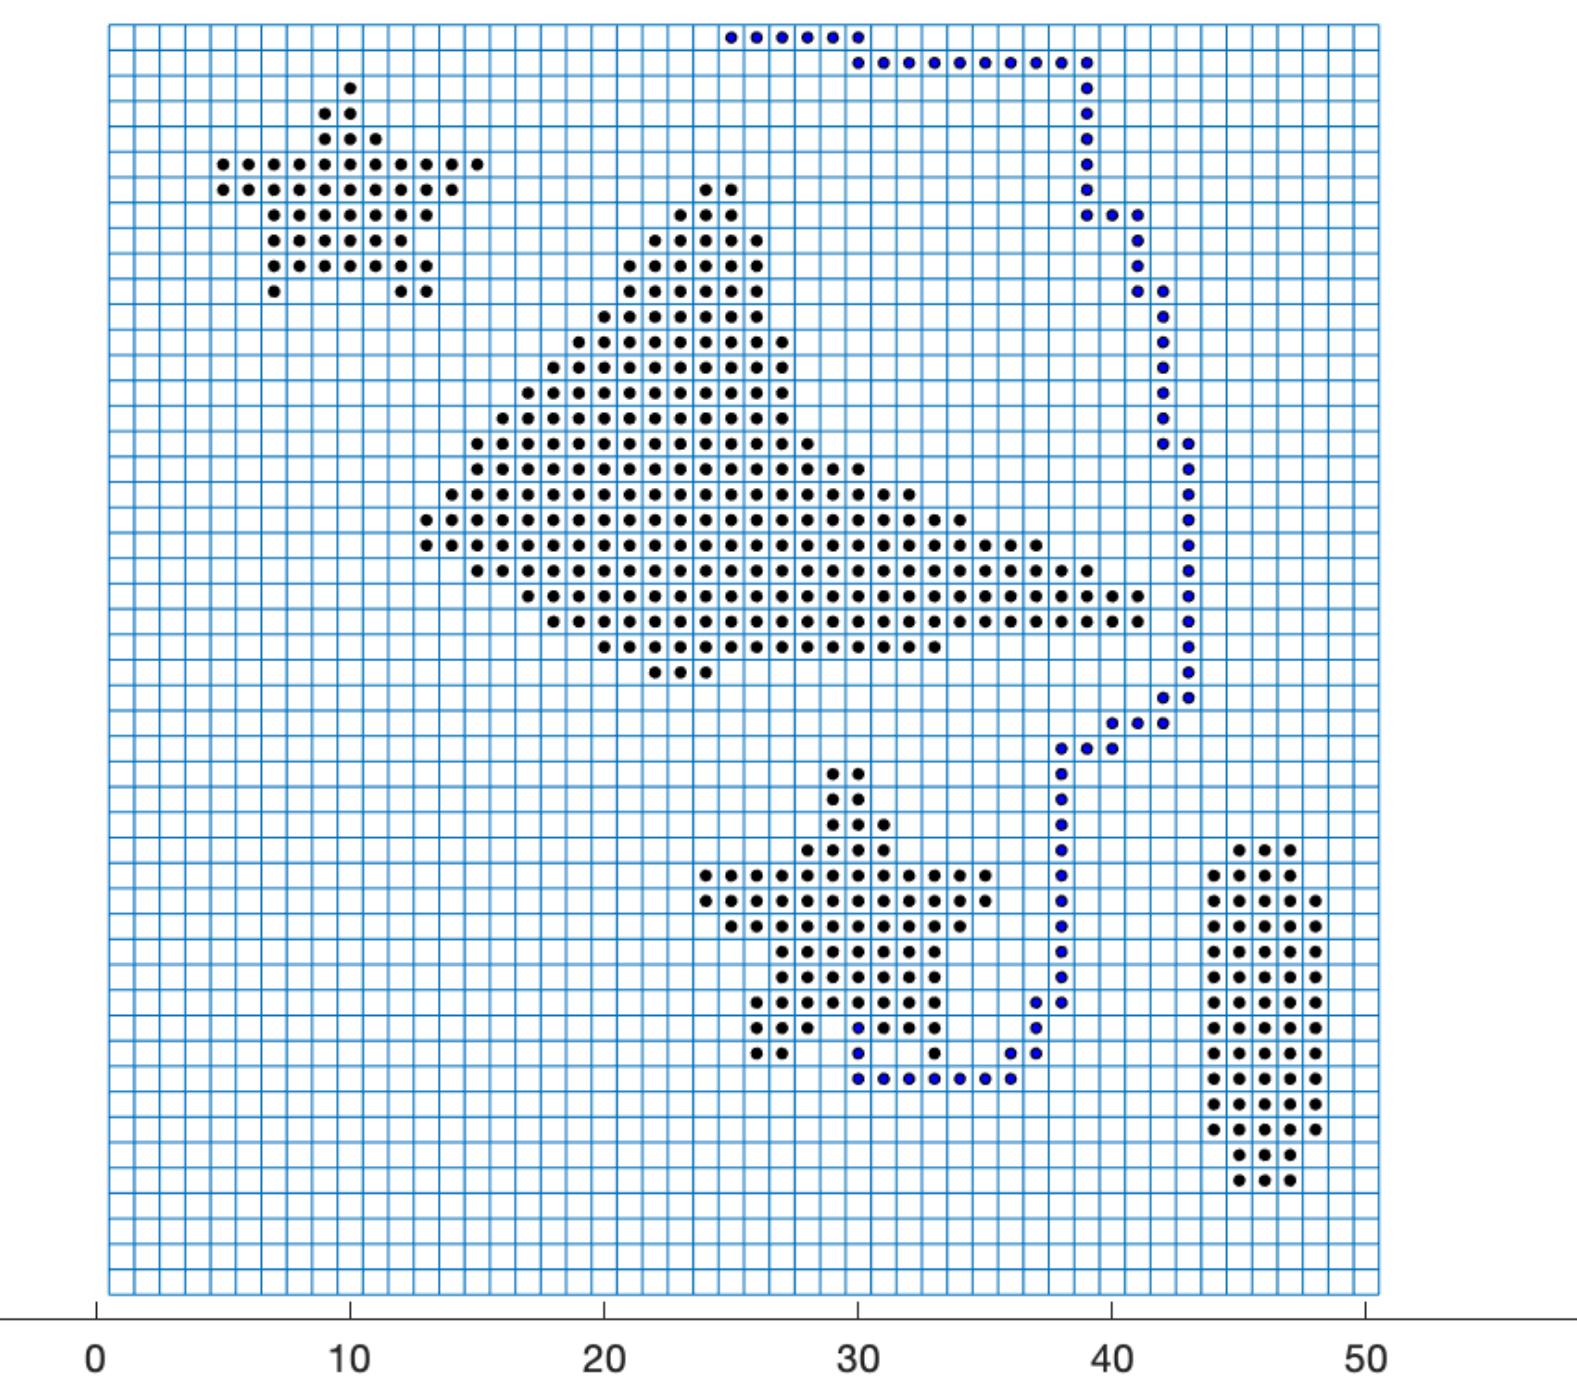


- **State:** The state of the robot: position, moving speed, heading direction,
- **Action:** Moving left, right, forward, backward, accelerating, turning,
- **Reward:** a positive value for reaching the goal, a negative value for reaching a forbidden position, 0 value for others.

Example - Motion Planning

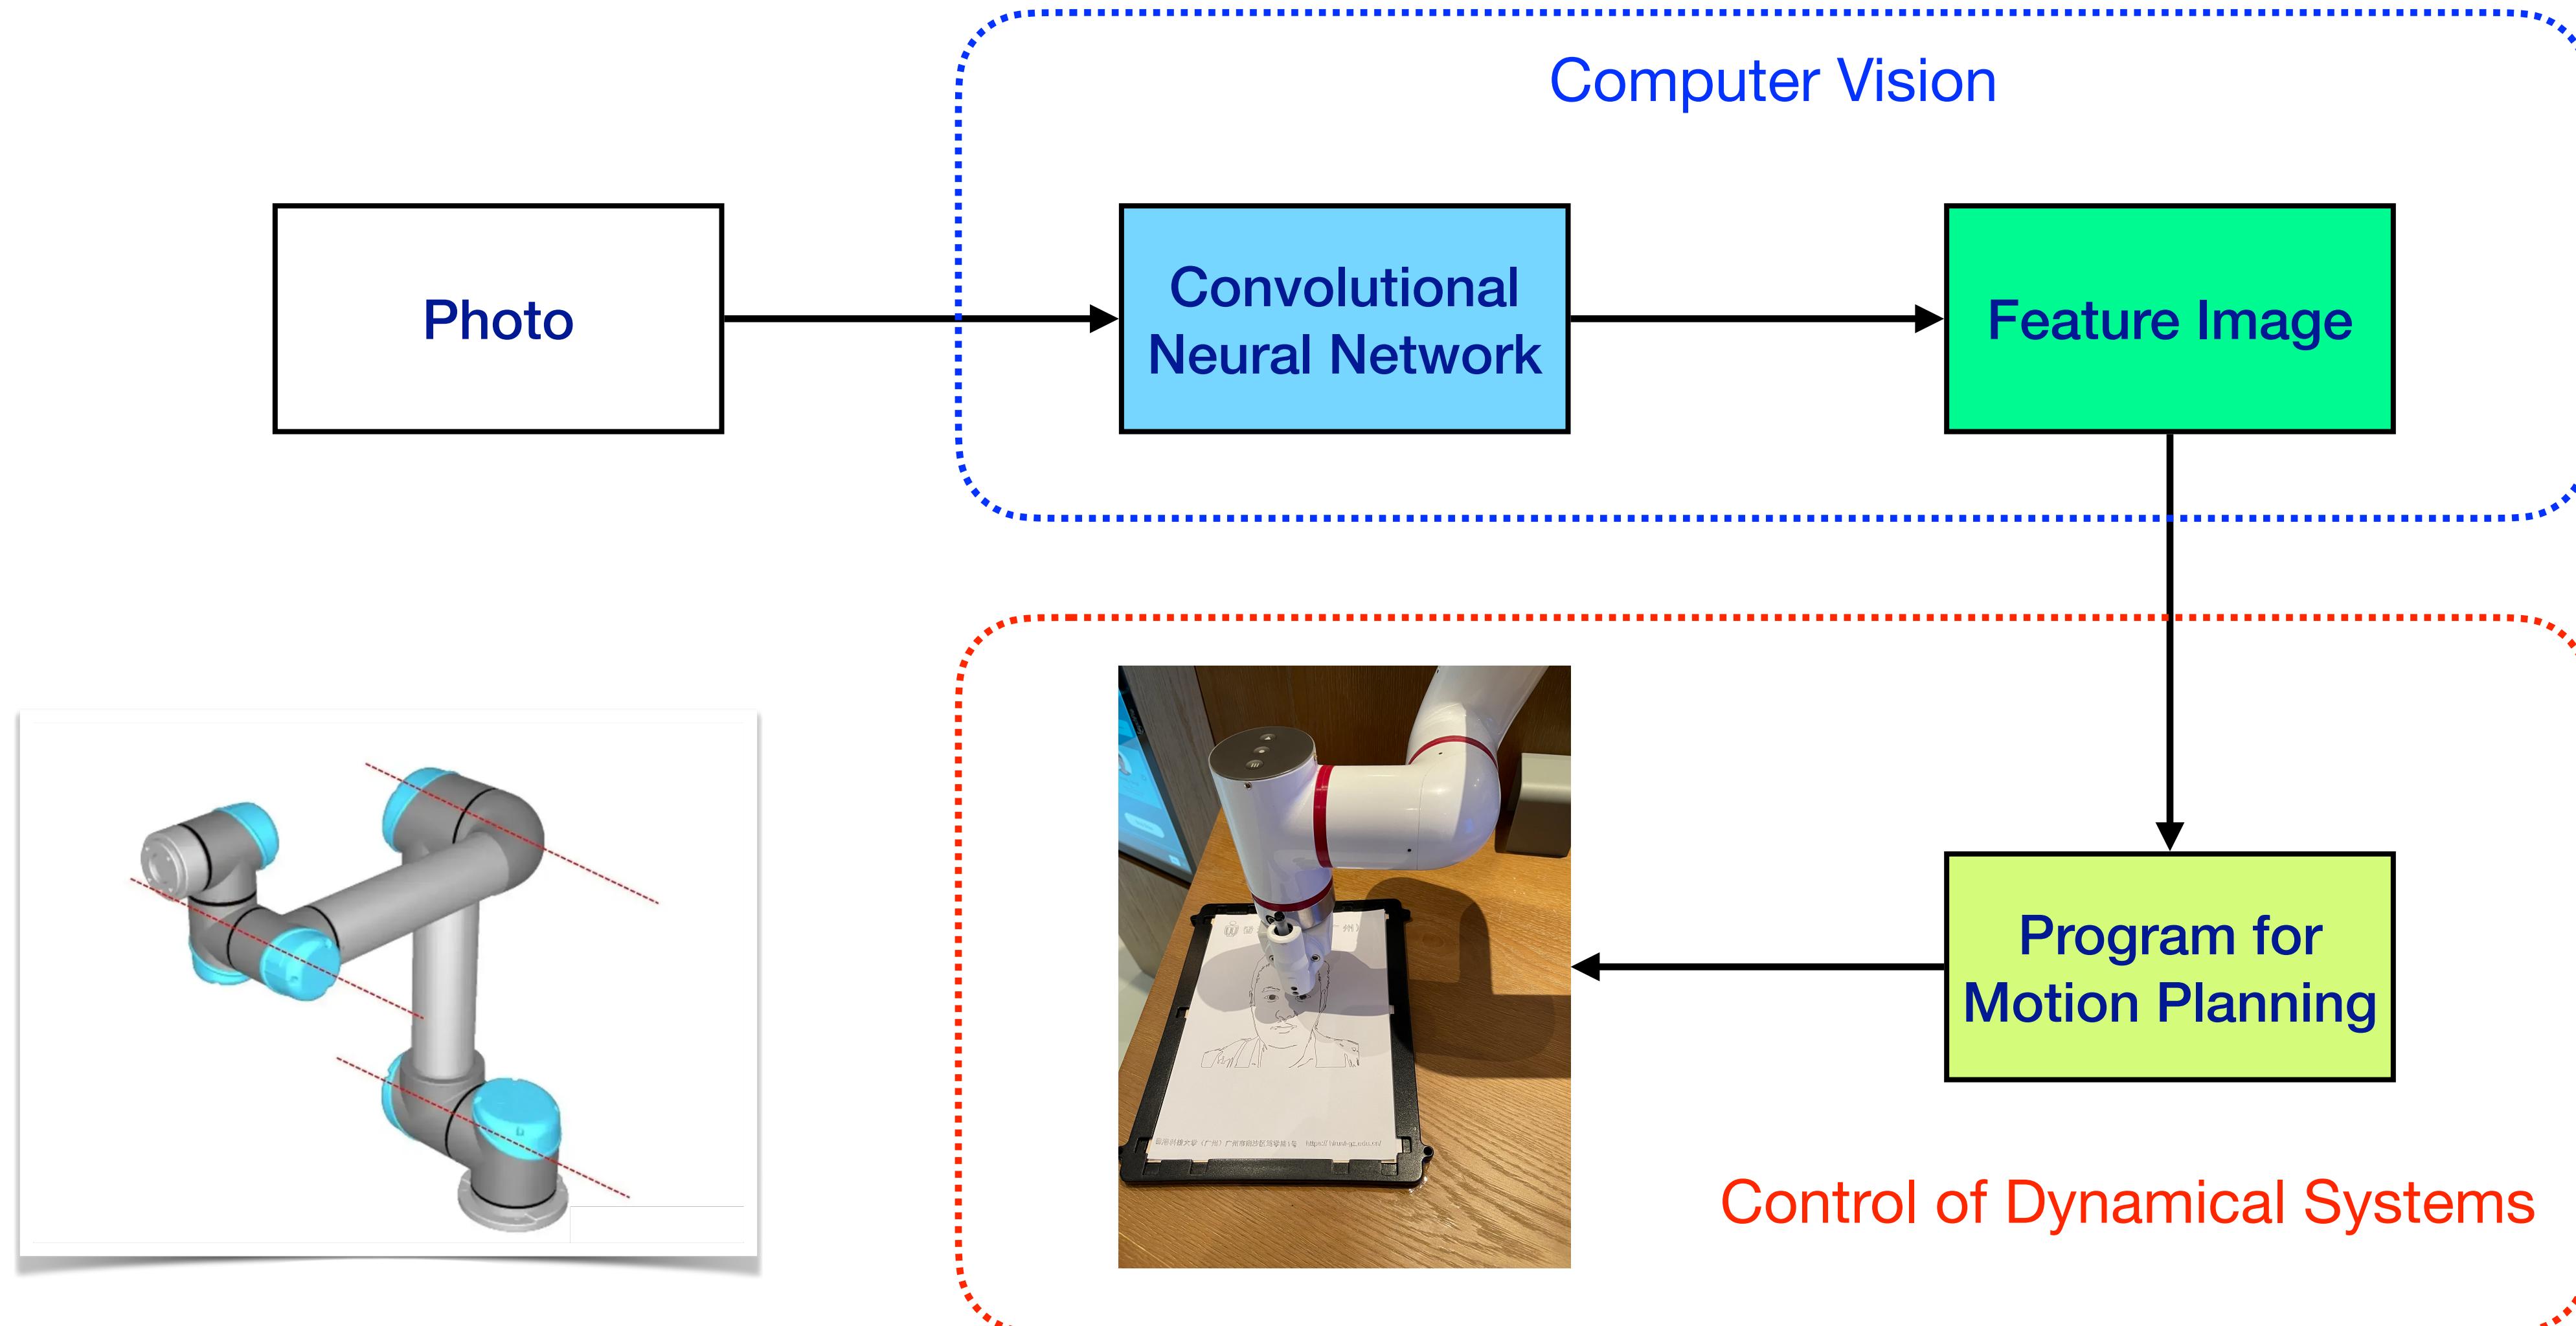


finite abstraction
→
(gridding the
continuous state
space)

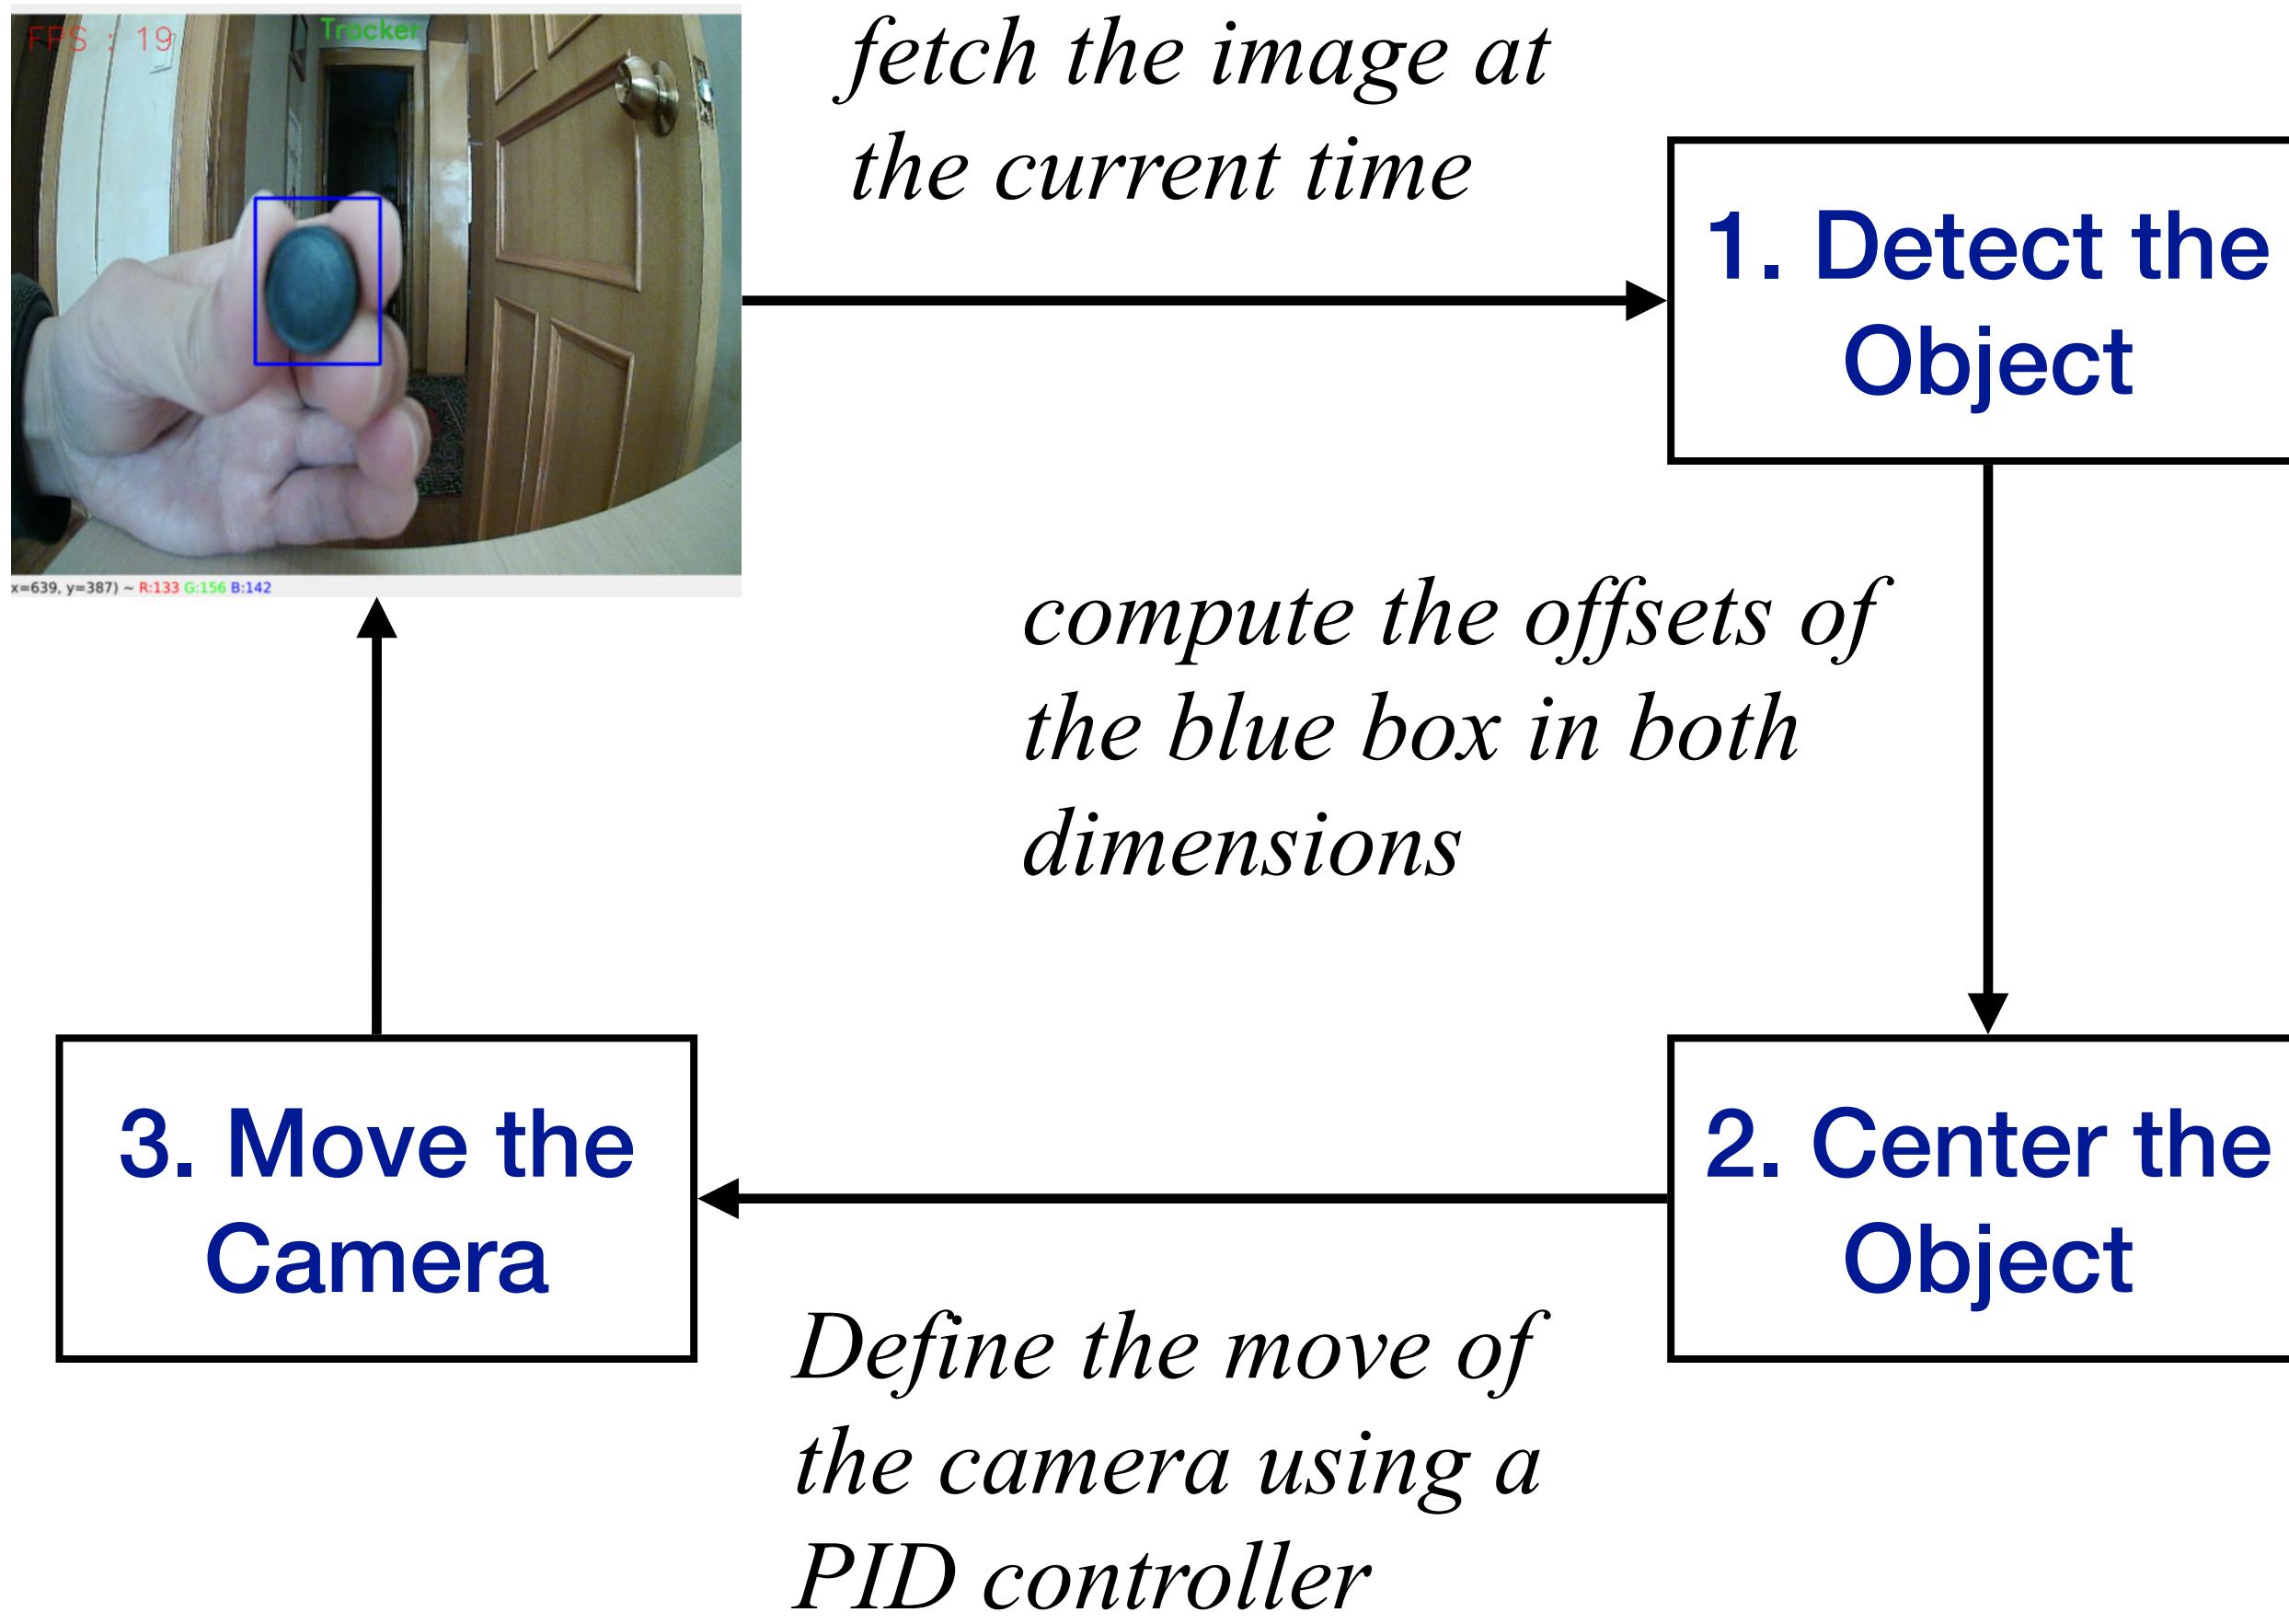


Walking from the start position (**red**) to the target position (**blue**) avoid reaching any obstacles. Four actions: **left**, **right**, **up** and **down**. After a reinforcement learning, the agent is going to tell the **best action** to choose in each position (grid).

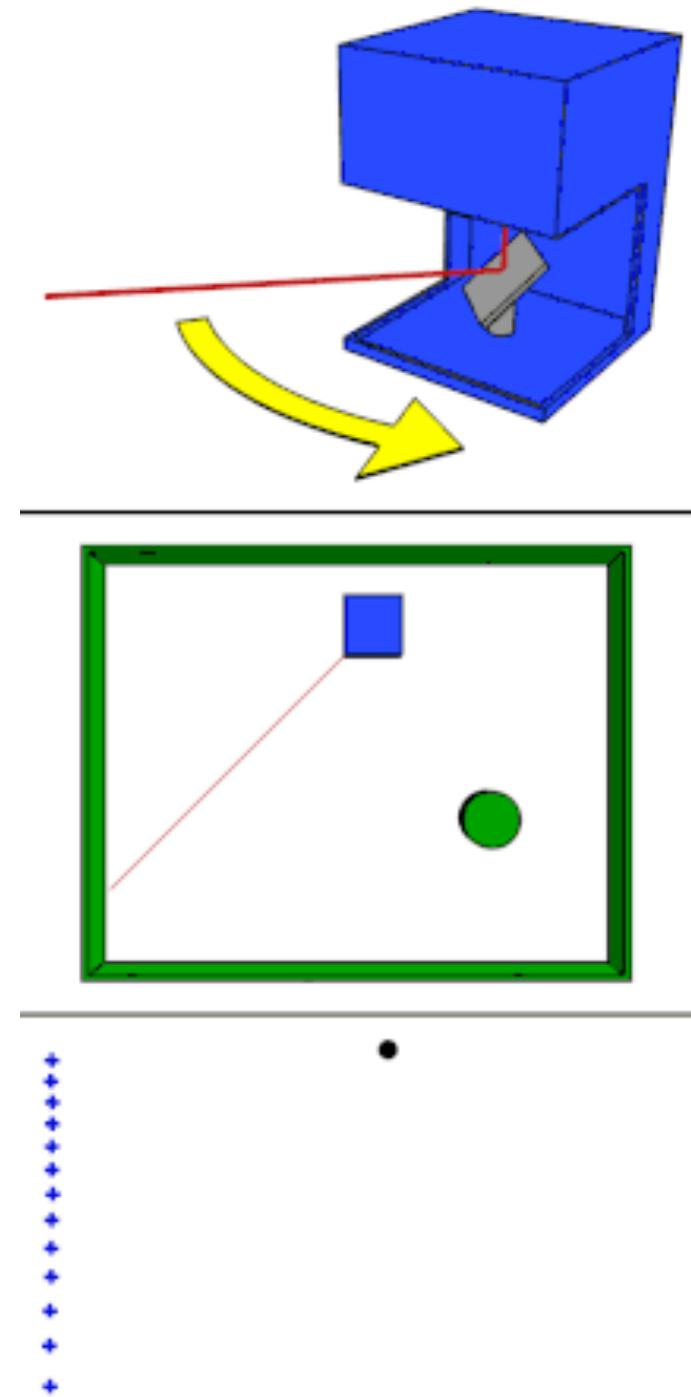
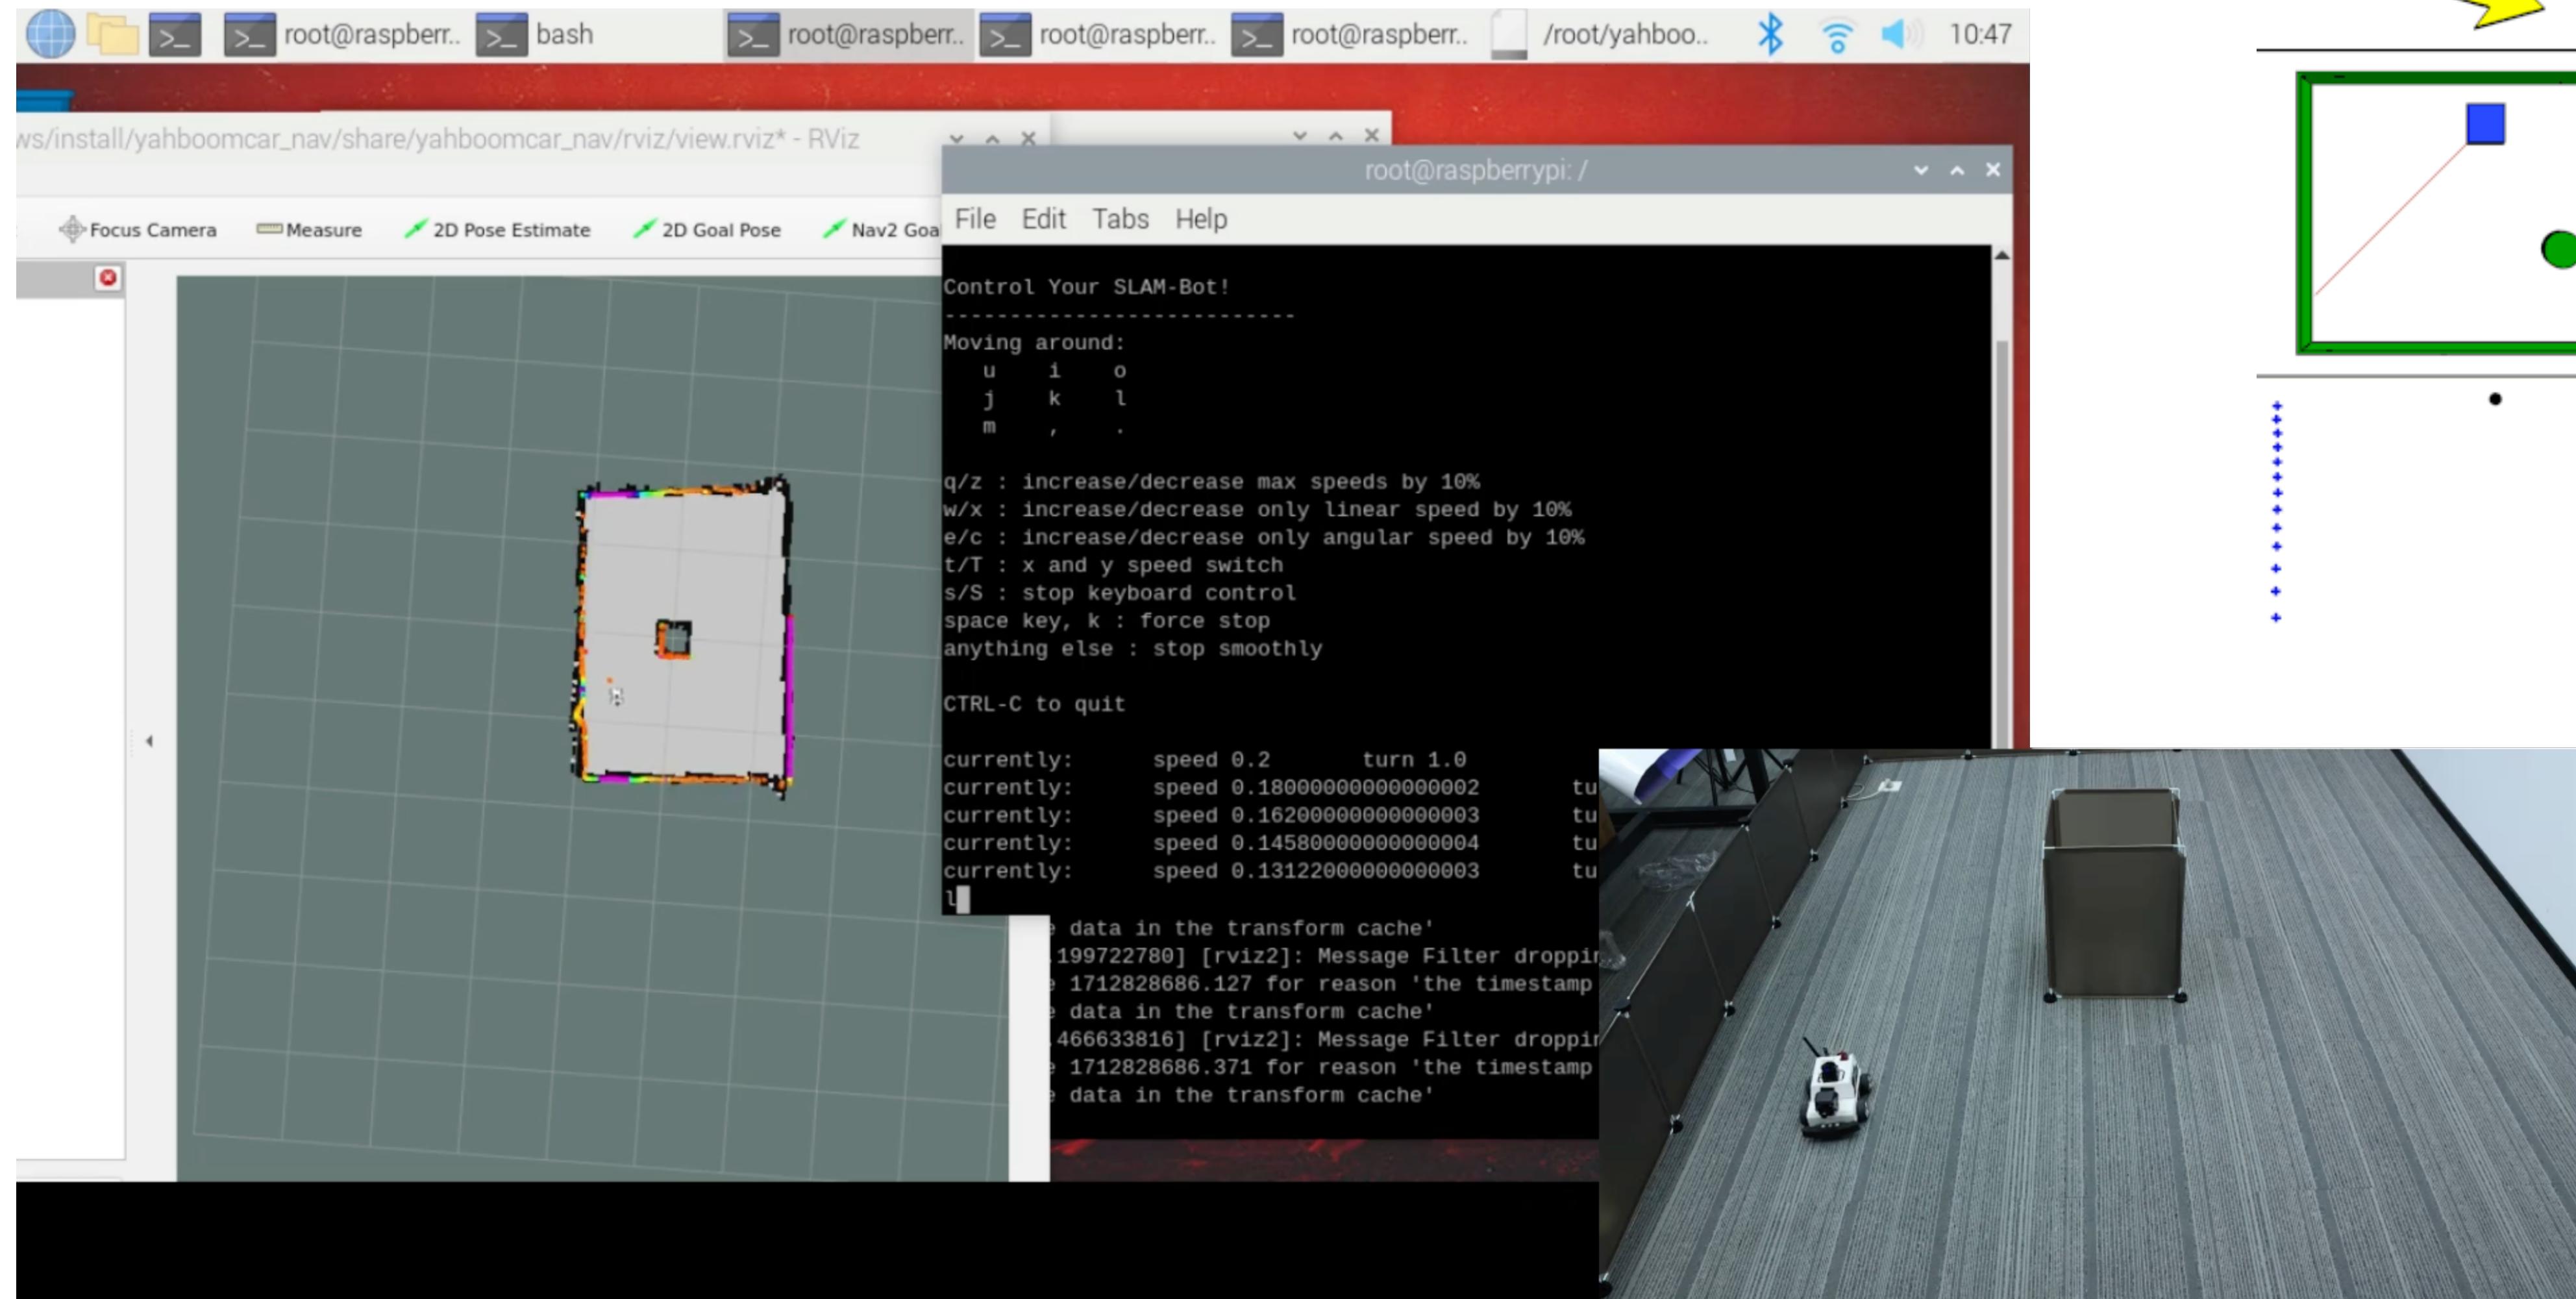
Motion Planning for Robot Arms



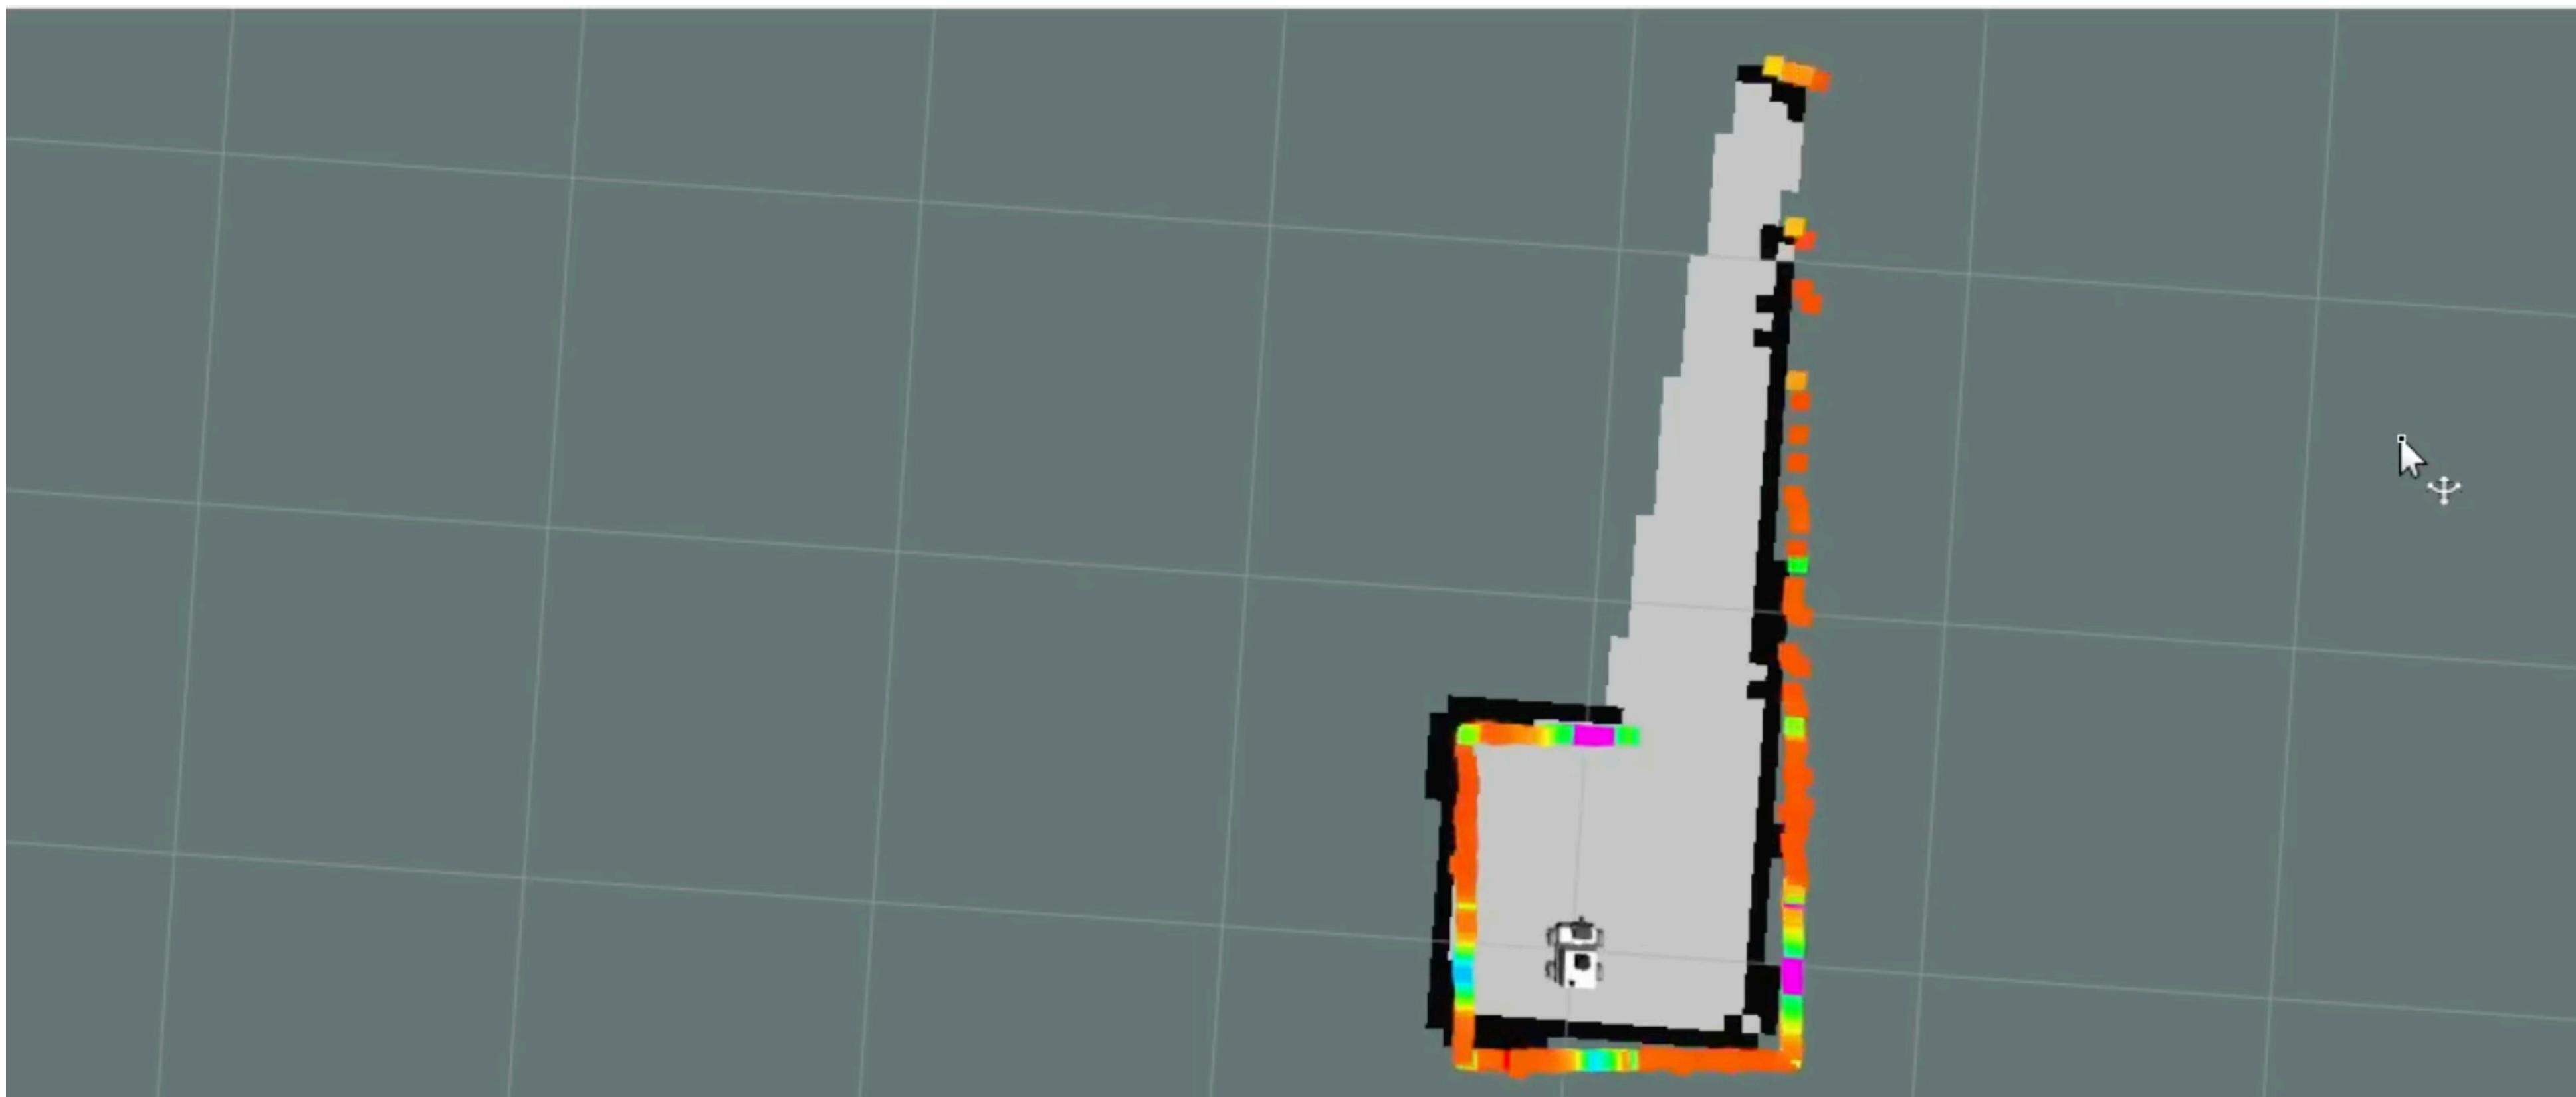
Object Tracking



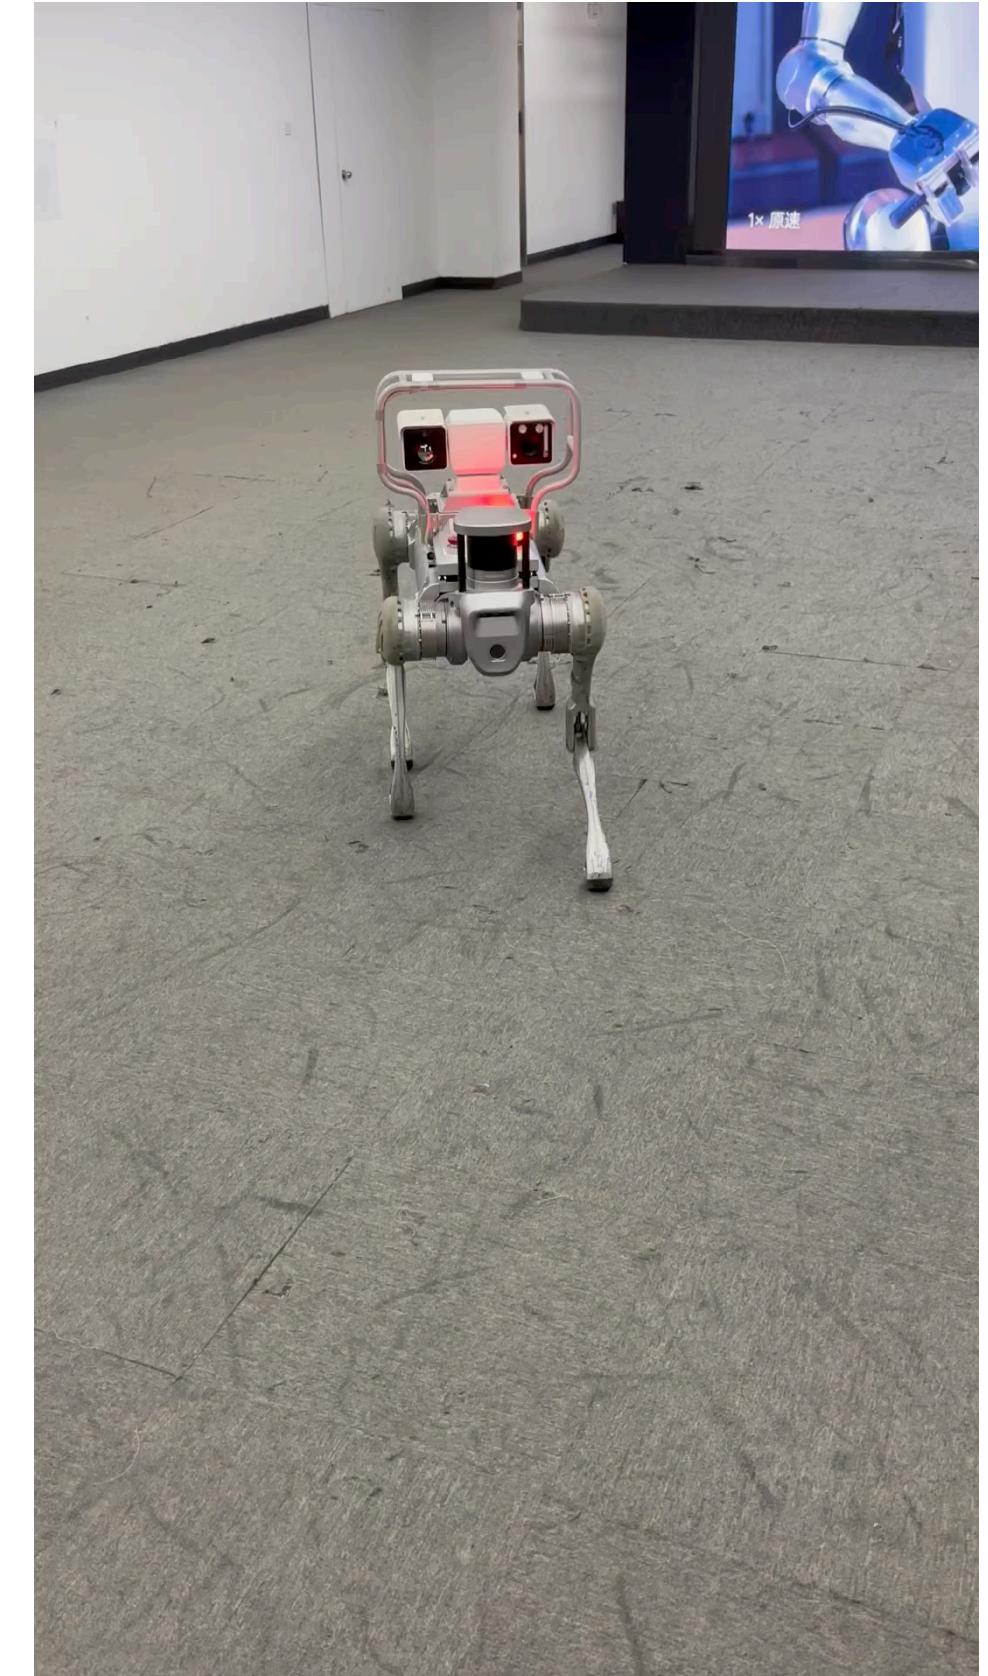
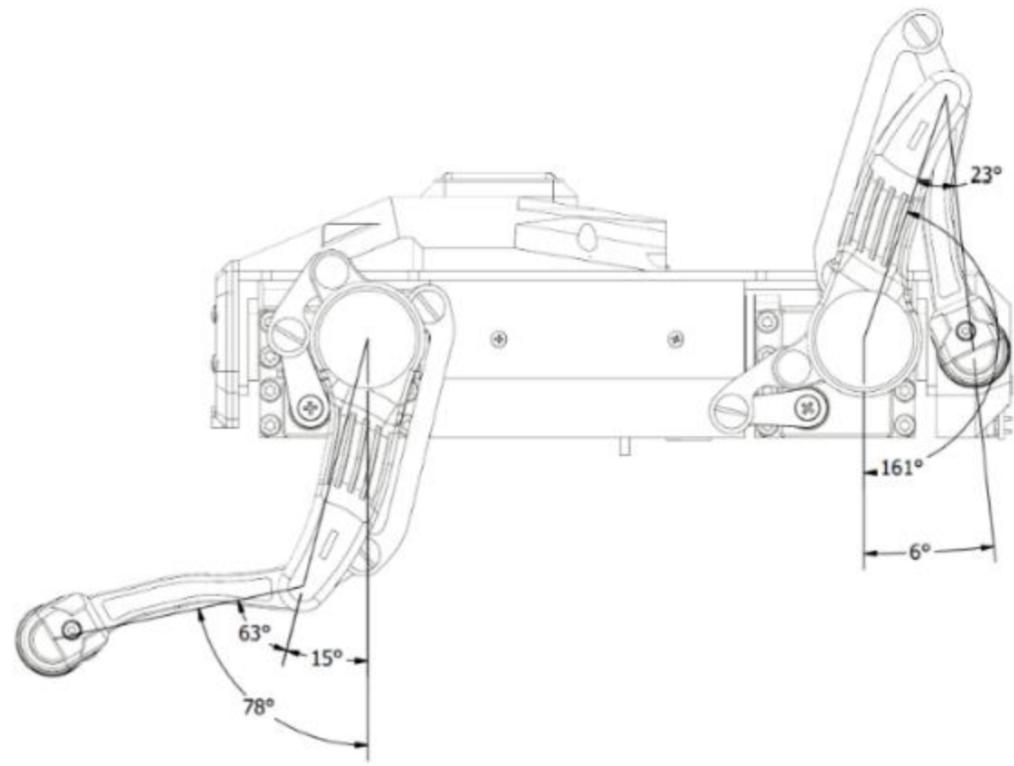
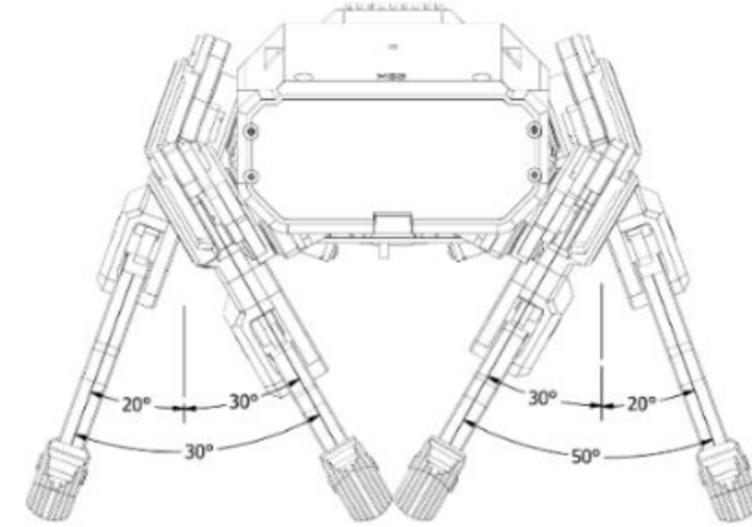
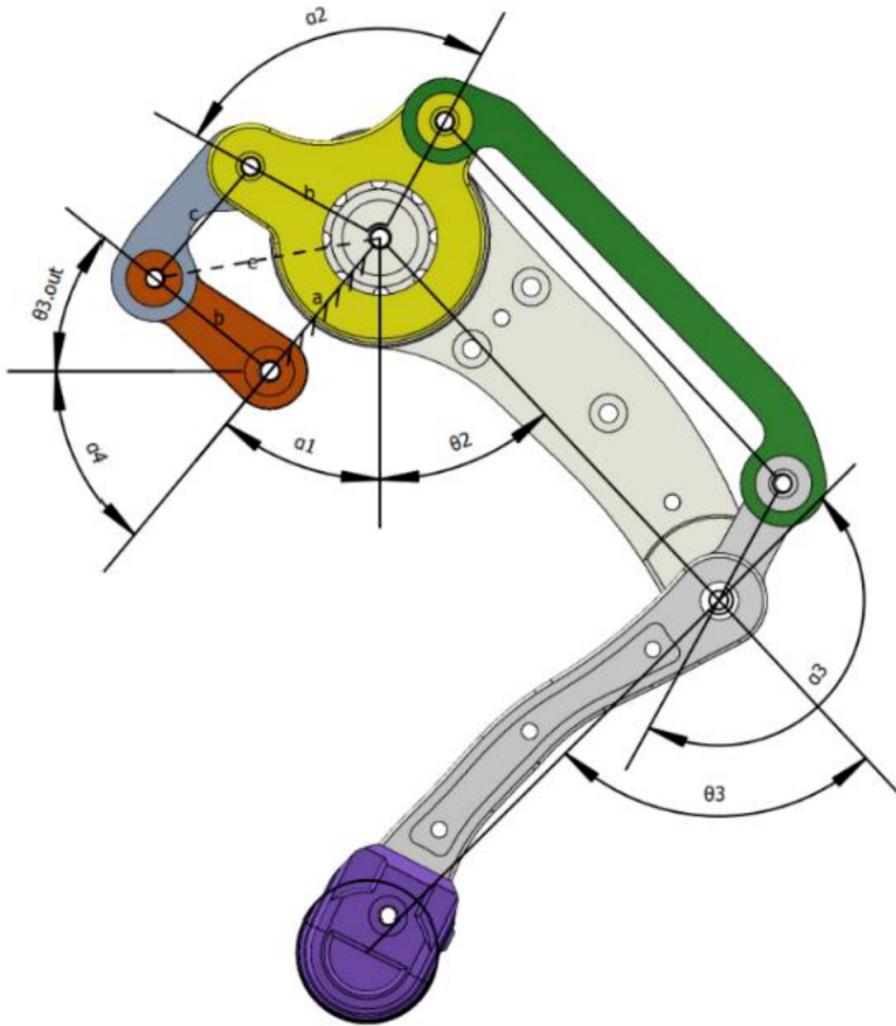
Building a Map from LiDAR Data



Navigation



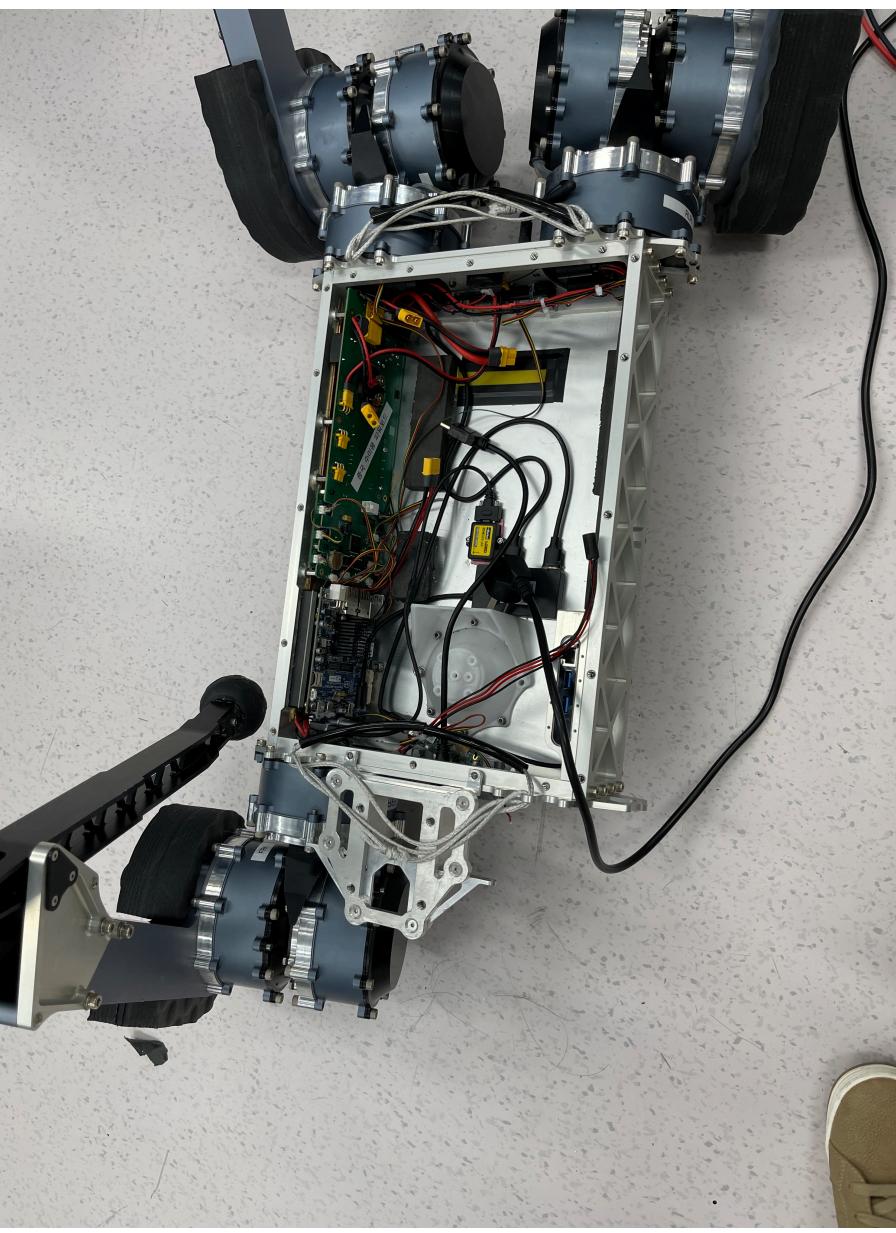
Control of Legged Robots



$$e^2 = a^2 + b^2 - 2ab \cos(\pi - \alpha_4 - \theta_{3.out})$$

$$\theta_3 = \theta_2 + \cos^{-1}\left(\frac{a^2 + e^2 - b^2}{2ae}\right) + \cos^{-1}\left(\frac{b^2 + e^2 - c^2}{2be}\right) + \alpha_1 + \alpha_2 - \alpha_3$$

<https://www.unitree.com/b2>



[HKUST-GZ]

Most of the ML tasks require a partially defined model which is often called a **parametric model**.

The task of training the model becomes an **optimization problem** which requires to find the unknown parameters that minimize model's prediction error.

We are going to learn how to solve such optimization problems.